

An Overview of C++11/14



Leor Zolman
BD Software
www.bdsoft.com
leor@bdsoft.com



September 8th, 2014

Agenda

- C++ timeline
- Goals for the new C++
- Part I. Simpler language changes
- Part II. New facilities for class design
- Part III. Larger new language features
 - Initialization-related improvements
 - Rvalue references, move semantics and perfect forwarding
 - Lambdas
- Most new language features are at least mentioned

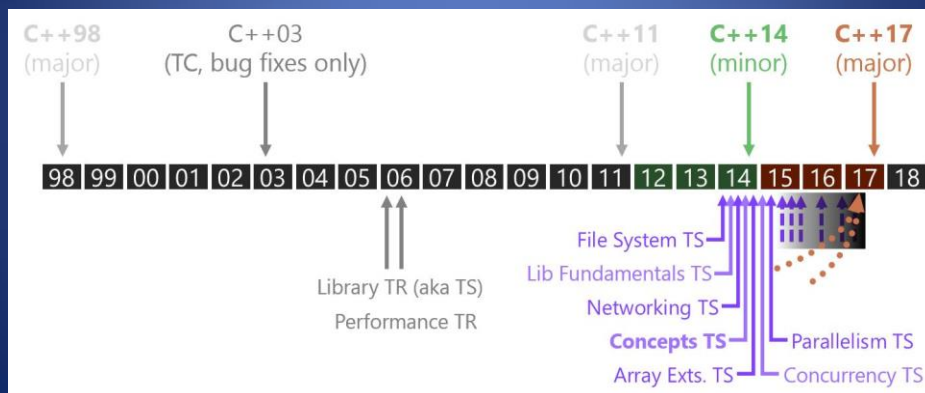
About the Code Examples

- When practical, I show specific problems/issues in Old C++ and then introduce the C++11/14 solutions
- Examples are not all 100% self-contained
 - Read the code *as if* the requisite `#includes`, `usings`, `std::s` etc. were there ☺

3

ISO C++ Timeline

(From Herb Sutter's Fall '03 ISO C++ Standards Meeting Trip Report, 10/3/13)



4

Goals for C++11

- Make C++ easier to teach, learn and use
- Maintain backward-compatibility
- Improve performance
- Strengthen library-building facilities
- Interface more smoothly with modern hardware

5

"The pieces just fit together better than they used to and I find a higher-level style of programming more natural than before and as efficient as ever."

-Bjarne Stroustrup [from his C++11 FAQ]

6

Part I: The Simpler Core Language Features

- `auto`, `decltype`, trailing return type
- `nullptr`
- Range `for`
- `>>` in template specializations
- `static_assert`
- `extern template`
- `noexcept`
- Variadic templates (OK, maybe *not* so simple)
- `constexpr`, template alias, and more...

7

Problem: Wordy declarations

```
// findNull: Given a container of pointers, return an  
// iterator to the first null pointer (or the end  
// iterator if none is found)
```

```
template<typename Cont>  
typename Cont::const_iterator findNull(const Cont &c)  
{  
    typename Cont::const_iterator it;  
    for (it = c.begin(); it != c.end(); ++it)  
        if (*it == 0)  
            break;  
  
    return it;  
}
```

8

Using findNull in Old C++

```
int main()
{
    int a = 1000, b = 2000, c = 3000;
    vector<int *> vpi;
    vpi.push_back(&a);
    vpi.push_back(&b);
    vpi.push_back(&c);
    vpi.push_back(0);

    vector<int *>::const_iterator cit = findNull(vpi);
    if (cit == vpi.end())
        cout << "no null pointers in vpi" << endl;
    else
    {
        vector<int *>::difference_type pos = cit - vpi.begin();
        cout << "null pointer found at pos. " << pos << endl;
    }
}
```

9

Using findNull in C++11

```
int main()
{
    int a = 1000, b = 2000, c = 3000;
    vector<int *> vpi { &a, &b, &c, 0 };

    auto cit = findNull(vpi);

    if (cit == vpi.end())
        cout << "no null pointers in vpi" << endl;
    else
    {
        auto pos = cit - vpi.begin();
        cout << "null pointer found in position " <<
            pos << endl;
    }
}
```

10

Problem: What's the Return Type?

- Sometimes a return type simply cannot be expressed in the usual manner:

```
// Function template to return product of two
// values of unknown types:

template<typename T, typename U>
??? product(const T &t, const U &u)
{
    return t * u;
}
```

11

decltype and Trailing Return Type

- In this case, a combination of `auto`, `decltype` and *trailing return type* provide the only solution for C++11:

```
// Function template to return product of two
// values of unknown types:

template<typename T, typename U>
auto product(const T &t, const U &u) -> decltype (t * u)
{
    return t * u;
}
```

12

findNull in C++11 (First Cut)

```
// findNull: Given a container of pointers, return an
// iterator to the first null pointer (or the end
// iterator if none is found)
```

```
template<typename Cont>
auto findNull(const Cont &c) -> decltype(c.begin())
{
    auto it = c.begin();
    for (; it != c.end(); ++it)
        if (*it == 0)
            break;
    return it;
}
```

13

Non-Member begin/end

- New forms of `begin()` and `end()` even work for native arrays, hence are more generalized

```
bool strLenGT4(const char *s) { return strlen(s) > 4; }
```

```
int main()
{
    // Applied to STL container:
    vector<int> v {-5, -19, 3, 10, 15, 20, 100};
    auto first3 = find(begin(v), end(v), 3);

    if (first3 != end(v))
        cout << "First 3 in v = " << *first3 << endl;
    // Applied to native array:
    const char *names[] {"Huey", "Dewey", "Louie"};
    auto firstGT4 = find_if( begin(names), end(names),
                           strLenGT4);

    if (firstGT4 != end(names))
        cout << "First long name: " << *firstGT4 << endl;
}
```

14

Non-Member begin/end Variations in C++14

- Return `const_iterator`s:
 - `cbegin/cend`
- Return `reverse_iterator`s:
 - `rbegin/rend`
- Return `const_reverse_iterator`s:
 - `crbegin/crend`

```
template <typename Container>
void process_container(Container &c)    // Note: no const
{
    typename C::const_iterator ci = begin(c); // C++11

    auto ci2 = cbegin(c);                // C++14

    . . .
}
```

15

Problem: Null Pointers

- In Old C++, the concept of “null pointers” can be a source of confusion and ambiguity
 - How is `NULL` defined?
 - Does `0` refer to an `int` or a pointer?

```
void f(long) { cout << "f(long)\n"; }
void f(char *) { cout << "f(char *)\n"; }

int main()
{
    f(0L);           // calls f(long)
    f(0);            // ERROR: ambiguous!
    f(static_cast<char *>(0)); // Oh, OK...
}
```

16

nullptr

- Using `nullptr` instead of `0` disambiguates:

```
void f(long) { cout << "f(long)\n"; }
void f(char *) { cout << "f(char *)\n";}

int main()
{
    f(0L);           // as before, calls f(long)
    f(nullptr);      // fine, calls f(char *)
    f(0);            // still ambiguous
}
```

17

findNull in C++11

(Final C++11 version)

```
template<typename Cont>
auto findNull(const Cont &c) ->
    decltype(begin(c))
{
    auto it = begin(c);
    for (; it != end(c); ++it)
        if (*it == nullptr)
            break;

    return it;
}
```

18

Generalized Function Return Type Deduction in C++14

- C++14 allows return type to be *deduced* from the return expression(s) used:

```
template<typename Cont>
auto findNull(const Cont &c)    // don't need decltype!
{
    auto it = begin(c);
    for (; it != end(c); ++it)
        if (*it == nullptr)
            break;

    return it;                  // return type deduced HERE
}
```

19

C++14: auto vs. decltype(auto)

- There are actually two approaches to function return type deduction in C++14
 - Functions declared to return **auto** (or “decorated” **auto**)
 - Employs *template* type deduction rules
 - Discards references, **const**, **volatile** from return expression's type (may add it back when **auto** is decorated)
 - Function declared to return **decltype(auto)**
 - No decoration permitted
 - Employs **decltype** type deduction
 - Expression's actual type is the return type

20

Iterating Over an Array or Container in Old C++

```
int main()
{
    int ai[] = { 10, 20, 100, 200, -500, 999, 333 };
    const int size = sizeof ai / sizeof *ai;    // A pain

    for (int i = 0; i < size; ++i)
        cout << ai[i] << " ";
    cout << endl;

    list<int> li ( ai, ai + size);
    // Note opportunities for typos, off-by-1's, etc.
    for (list<int>::iterator it =
        li.begin(); it != li.end(); ++it)
        *it += 100000;

    // Same rigmarole here:
    for (list<int>::const_iterator it =
        li.begin(); it != li.end(); ++it)
        cout << *it << " ";
}
```

21

Improvement: Range-Based for Loop

```
int main()
{
    int ai[] { 10, 20, 100, 200, -500, 999, 333 };

    for (auto i : ai)    // Don't need size
        cout << i << " ";
    cout << endl;

    list<int> li (begin(ai), end(ai));
    for (auto &elt : li)  // Using a ref allows modifying
        elt += 10000;

    for (const auto &elt : li) // Note: const & not nec-
        cout << elt << " ";    // essarily better for ints
    cout << endl;

    for (auto i : { 100, 200, 300, 400 })
        cout << i << " ";
}
```

22

The “>> Problem”

- Old C++ requires spaces between consecutive closing angle-brackets of nested template specializations:

```
map<string, vector<string> > dictionary;
```

- C++11 permits you to omit the space:

```
map<string, vector<string>> dictionary;
```

- That's one less *gotcha*

23

Compile-Time Assertions: `static_assert`

- The C library contributed the venerable `assert` macro for expressing run-time invariants:

```
int *pi = ...;  
assert (pi != NULL);
```

- C++11 provides direct language support for *compile-time* invariant validation and diagnosis:

```
static_assert(condition, "message");
```

- Conditions may only be formulated from *constant* (compile-time determined) expressions

24

static_assert

```
static_assert(sizeof(int) >= 4,
    "This app requires ints to be at least 32 bits.");

template<typename R, typename E>
R safe_cast(const E &e)
{
    static_assert(sizeof(R) >= sizeof(E),
        "Possibly unsafe cast attempt.");
    return static_cast<R>(e);
}

int main()
{
    long lval = 50;
    int ival = safe_cast<int>(lval); // OK iff long & int
                                    // are same size
    char cval = safe_cast<char>(lval); // Compile error!
}
```

25

Problem: Object File Code Bloat From Templates

- The industry has settled on the “template inclusion model”
 - Templates fully defined in header files
 - Each translation unit (module) `#includes` the header: all templates are instantiated in *each* module which uses them
 - At link time, all but one instance of each redundant instantiated function *is discarded*

26

The Failed Solution: `export`

- Old C++ introduced the `export` keyword
- The idea was to support *separately compiled templates*
- But even when implemented (AFAIK only EDG accomplished this), *it didn't really improve productivity*
 - Templates are just too complicated
 - ...due to two-phase translation

27

The C++11 Solution: `extern template`

- Declare a class template specialization `extern` and the compiler will not instantiate the template's functions in that module:

```
#include <vector>
#include <widget>
extern template class vector<widget>;
```

- For `vector<widget>`, the *class definition* is generated if needed (for syntax checking) but member functions are not instantiated
- Then, in just *one* (`.cpp`) module, *explicitly instantiate* the class template:

```
template class vector<widget>;
```

28

Problem: Dynamic Exception Specifications

- In Java, all exception specifications are *enforced*
- Old C++ functions specify exceptions they might throw via *dynamic exception specifications*...but callers need not acknowledge them!
- Plus, how can function *templates* possibly know what exceptions might be thrown?
- Thus the only dynamic exception specification used in the Old C++ standard library is the *empty* one:

```
template<typename T>
class MyContainer {
public:
    ...
    void swap(MyContainer &) throw();
    ...
}
```

29

The C++11 Way: `noexcept`

- Dynamic exception specifications (even empty ones) can impact performance
- C++11 *deprecates* dynamic exception specifications and introduces the `noexcept` keyword:

```
template<typename T>
class MyContainer {
public:
    ...
    void swap(MyContainer &) noexcept;
    ...
}
```

- If an exception tries to escape from a `noexcept` function, the program immediately terminates.

30

Conditional noexcept

- `noexcept` clauses can be conditional on the “`noexcept`” status of sub-operations

```
// From the g++ standard library's implementation
// of std::pair (simplified):
template<class T1, class T2>
class pair
{
    // Defines data members 'first', 'second', etc.

    void swap(pair& __p)
        noexcept(noexcept(swap(first, __p.first))
            && noexcept(swap(second, __p.second)))
    {
        using std::swap;
        swap(first, __p.first);
        swap(second, __p.second);
    }
}
```

31

Problem: How Do You Write a Function to Average N Values?

- You can use C variadic functions:


```
int averInt(int count, ...);
double averDouble(int count, ...);
```

 - Must write one for each type required
 - Must provide the argument count as 1st arg
 - Type safety? Fuggedaboutit...
- Can't use C++ default arguments
 - Because we can't know the # of actual args
- Could use overloading and templates
 - That's ugly too

32

Variadic Templates

```
// To get an average, we 1st need a way to get a sum...
template<typename T> // ordinary function template
T sum(T n)           // for the "terminal" case
{
    return n;
}

// variadic function template:
template<typename T, typename... Args>
T sum(T n, Args... rest) // "parameter packs"
{
    return n + sum(rest...);
}

int main() {
    cout << sum(1,2,3,4,5,6,7);
    cout << sum(3.14, 2.718, 2.23606);
};
```

33

A Subtle Problem With sum

- Note that if `sum` is called with a mixture of different argument types, the wrong return type may result:

```
template<typename T, typename... Args>
T sum(T n, Args... rest) // T is the FIRST type
{
    return n + sum(rest...);
} // sum(1, 2.3) yields 3 (instead of 3.3)
```

- To fix this, first we need `auto`, trailing return type, etc.:

```
template<typename T, typename... Args>
auto sum(T n, Args... rest) -> decltype(n + sum(rest...))
{
    return n + sum(rest...);
}
```

- Unfortunately, *that* doesn't compile ☹ ☹

34

An Even More Subtle Problem

- The recursive reference to `sum` in the `decltype` expression is illegal because the compiler doesn't know the full type of `sum` (haven't reached the end of its header line yet...)
- A surprising exception to this restriction: when the function name is of a *member* function (I kid you not) it is OK if it is still "incomplete"
- So to make it all work, just make the `sum` function templates be `static` functions of a struct!

35

Finally, a Working sum

```
struct Sum
{
    template<typename T>
    static T sum(T n)
    {
        return n;
    }

    template<typename T, typename... Args>
    static auto sum(T n, Args... rest) ->
        decltype(n + sum(rest...))
    {
        // NOW the recursion is OK...
        return n + sum(rest...);
    }
};
```

36

And, At Last: Average

- Another variadic function template can leverage the `sum()` templates, and variadic `sizeof...` operator, to give us average:

```
template<typename... Args>
auto avg(Args... args) ->
    decltype(Sum::sum(args...))
{
    return Sum::sum(args...) / (sizeof... (args));
}

cout << avg(2.2, 3.3, 4.4) << endl; // works!
cout << avg(2, 3.3, 4L) << endl;    // works too!
```

37

constexpr

- Enables compile-time evaluation of functions (including operators and constructors) when expressed in terms of *constant* expressions

```
const int SIZE = 100;           // simple constant

template<typename T>
constexpr T square(const T&x)  // constexpr function--
{                               // serves "double duty"
    return x * x;
}

Widget warray[SIZE];           // OK, as always
Widget warray2[square(SIZE)];  // OK (compile-time)

int val = 50;                  // NOT constant expr.
int val_squared = square(val); // OK, run-time expr.
Widget warray3[square(val)];    // ERROR (not const.) 38
```

C++11 constexpr vs. C++14's

- Bodies of C++11 `constexpr` functions are essentially limited to a single return expression
- No other control structures allowed
 - No `if...else`
 - But `?:` can serve as a substitute
 - No loops
 - But recursion is supported
- C++14 relaxes most of C++11's restrictions
 - “Roughly: a `constexpr` function can contain anything that does not have side effects outside the function.” [--B. Stroustrup]
 - Still forbidden: `goto`, `try` blocks, calling non-`constexpr` functions, a few other minor things ³⁹

Template Alias

- The “template typedef” idea, w/clearer syntax:

```
template<typename T>
using setGT = std::set<T, std::greater<T>>;

setGT<double> sgtd { 1.1, 8.7, -5.4 };
// As if writing:
// std::set<double, std::greater<double>> sgtd {...
```

- `using` aliases also make a “better typedef”:

```
typedef void (*voidfunc)();      // Old way
using voidfunc = void (*)();     // New way
```

Some String-Related Features

- Unicode string literals
 - UTF-8: `u8"This text is UTF-8"`
 - UTF-16: `u"This text is UTF-16"`
 - UTF-32: `U"This text is UTF-32"`
- Raw string literals
 - Can be clearer than lots of escaping:

```
string s = "backslash: \\\"\\\", single quote: '\\\"';  
string t = R"(backslash: \"\", single quote: '\"')";  
// Both strings initialized to:  
//      backslash: "\", single quote: '"  
  
string u = R"xyz(And here's how to get )" in!)xyz";
```

41

Inline Namespaces

- Facilitates versioning
- Names in an `inline` sub-namespace are implicitly “hoisted” into the enclosing (parent) namespace

```
namespace bdsoft {  
    namespace v1_0          // old version  
    { ... }  
    inline namespace v2_0   // Current/default version  
    {  
        ... // Everything in here behaves as if  
              // declared directly in namespace bdsoft  
    }  
    namespace v3_0          // Experimental version  
    { ... }  
}
```

42

Attributes

- Replaces #pragmas, __attribute__, __declspec, etc.
- E.g., `[[noreturn]]` to help compilers detect errors
- New in C++14: `[[deprecated]]`
 - Compiler issues warning if labeled entity used
 - Can be used for: functions, classes, typedefs, enums, variables, non-static data members and template *complete specializations*

43

More Language Features

- Scoped `enums` (a.k.a. *enum classes*)
 - Enumerators don't "leak" into surrounding scopes
 - Fewer implicit conversions
 - Can specify the underlying (integral) type
 - This is true for old (un-scoped) enums as well
- `long long`
 - 64-bit (at least) ints
- `alignas` / `alignof`
 - Query/ force boundary alignment

44

Yet More Language Features

- Generalized Unions
 - E.g., members of unions are now allowed to have constructors, destructors and assignment
 - However, any user-defined ctor, dtor or copy op is treated as if it were declared `=delete` (and thus cannot be used with an object of the union type)
- Generalized PODs
 - E.g., “Standard Layout Types” (PODs) can now have constructors
 - C++98 POD types are now subdivided into: PODs, trivially copyable types, trivial types and standard-layout types

45

Yet More Language Features

- Garbage Collection ABI
 - Sets ground-rules for gc; specifies an ABI.
[Note: Actual gc is neither required nor supplied]
- User-defined Literals
 - Classes can define *literal operators* to convert from literals with a special suffix into objects of the class type, e.g.,
`Binary b = 11010101001011b;`

46

New for C++14

- Binary Literals
 - The new prefix `0b` (or `0B`) designates intrinsic binary literal values:
`auto fifteen = 0b1111;`
- Single-quote as Digit Separator:
`auto filler_word = 0xdead'beef;`
`auto aBillion = 1'000'000'000;`

47

Part II: Features Supporting Better Class Design

- Generated functions: `default` / `delete`
- Override control: `override` / `final`
- Delegating constructors
- Inheriting constructors
- Increased flexibility for in-class initializers
- Explicit conversion operators

48

Problem: How to Disable Copying?

- There are at least two Old C++ approaches to prevent objects from being copied:

- Make the copy operations private:

```
class RHC          // some resource-hogging class
{
...
private:
    RHC(const RHC &);
    RHC &operator=(const RHC &);
};
```

- Inherit privately from a base class that does it for you:

```
class RHC : private boost::noncopyable
{
...
};
```

- Both are problematic.

49

Now: =default, =delete

- These specifiers control function generation:

```
class T {
public:
    T() = default;
    T(const char *str) : s(str) {}
    T(const T&) = delete;
    T &operator=(const T&) = delete;
private:
    string s;
};

int main() {
    T t;           // Fine
    T t2("foo");  // Fine
    T t3(t2);      // Error!
    t = t2;        // Error!
}
```

50

=delete and Func. Overloading

- =delete also lets you restrict arg types:

```
void doSomething(int);
void doSometehing(long) = delete;
void doSomething(double) = delete;

int main()
{
    doSomething(10);           // Fine (direct)
    doSomething('c');          // Fine (conversion)
    doSomething(10L);          // ERROR
    doSomething(2.3);          // ERROR
    doSomething(2.3F);         // ERROR (!)
}
```

51
51

Problems With Overriding

- When limited to Old C++ syntax, the “overriding interface” is potentially misleading / error-prone:

```
class Base {
public:
    virtual void f(int);
    virtual void ff(int);
    virtual int g() const;
    void h(int);
};

class Derived : public Base {
public:
    void f(int);           // is this a virtual func.?
    virtual int g();       // meant to override Base::g?
    void ff(int);          // How to prevent further over.?
    void h(int);           // overrides Base::h? Or... ?
};
```

52

override / final

- C++11 (mostly) lets you say what you really mean:

```
class Base {
public:
    virtual void f(int);
    virtual void ff(int);
    virtual int g() const;
    void h(int); // final
};

class Derived : public Base {
public:
    void f(int) override; // Base::f MUST be virtual
    void ff(int) final;   // Prevents overriding!
    int g() override;    // Error!
    void h(int);         // SHOULD be an error...
                        // ... but no help here
};

// Note: These are "CONTEXTUAL" keywords! Cool!
```

53

final Classes

- An entire class can be declared final:

```
class Base {};
```



```
class Derived final : public Base { // Derived is final
    . . .
};
```



```
class Further_Derived : public Derived { // ERROR!
    . . .
};
```

54

Problem: Old C++ Ctors Can't Use the Class' Other Ctors

```
class FluxCapacitor
{
public:
    FluxCapacitor() : capacity(0), id(nextId++) {}
    FluxCapacitor(double c) : capacity(c),
                             id(nextId++) { validate(); }
    FluxCapacitor(complex<double> c) : capacity(c),
                                       id(nextId++) { validate(); }
    FluxCapacitor(const FluxCapacitor &f) :
        id(nextId++) {} // can you spot the
                        // silent logic error?
    // ...
private:
    complex<double> capacity;
    int id;
    static int nextId;
    void validate();
};
```

55

C++11 Delegating Constructors

- C++11 ctors may call other ctors (à la Java)

```
class FluxCapacitor
{
public:
    FluxCapacitor() : FluxCapacitor(0.0) {}
    FluxCapacitor(double c) :
        FluxCapacitor(complex<double>(c)) {}
    FluxCapacitor(const FluxCapacitor &f) :
        FluxCapacitor(f.capacity) {}
    FluxCapacitor(complex<double> c) :
        capacity(c), id(nextId++) { validate(); }
    // Note: Now harder to forget to set capacity
private:
    complex<double> capacity; // BUT... There's a
    int id;                  // subtle performance
    static int nextId;       // hit in this part-
    void validate();         // icular example...
};
```

56

Problem: Very Limited Data Member Initialization

- In old C++, *only* const static integral members could be initialized in-class

```
class FluxCapacitor
{
public:
    static const size_t num_cells = 50;    // OK
    FluxCapacitor(complex<double> c) :
        capacity(c), id(nextId++) {}
    FluxCapacitor() : id(nextId++) {}      // capacity??
    ...
private:
    int id;
    complex<double> capacity = 100;       // ERROR!
    static int nextId = 0;                 // ERROR!
    Cell FluxCells[num_cells];            // OK
};
```

57

C++11 In-Class Initializers

- Now, any non-static data member can be initialized in its declaration:

```
class FluxCapacitor
{
public:
    static const size_t num_cells = 50;    // still OK
    FluxCapacitor(complex<double> c) :
        capacity(c), id(nextId++) {}      // capacity c
    FluxCapacitor() : id(nextId++) {}      // capacity 100
    ...
private:
    int id;
    complex<double> capacity = 100;        // Now OK!
    static int nextId = 0;                 // Still illegal
    Cell FluxCells[num_cells];            // Still OK
};
```

58

Inheriting Constructors

- C++11 derived classes may “inherit” most ctors (just not the default ctor) from their base class(es):
 - Simply extends the old `using Base::name` syntax to ctors (arbitrarily excluded in C++98)
 - New ctors may still be added
 - Inherited ones may be redefined

```
class RedBlackFluxCapacitor : public FluxCapacitor
{
public:
    enum Color { red, black };
    using FluxCapacitor::FluxCapacitor;
    RedBlackFluxCapacitor(Color c) : color(c) {}
    void setColor(Color c) { color = c; }
private:
    Color color { red };    // Note: default value
};
```

59

Explicit Conversion Operators

- In Old C++, only constructors (one form of user-defined conversion) could be declared `explicit`
- *Operator* conversion functions (e.g., `operator long()`) could not
- C++11 remedies that, *but...*

```
class Rational {
public:
    // ...
    operator double() const;           // Iffy...
    explicit operator double() const;  // Better...
    double toDouble() const;           // Best?
private:
    long num, denom;
};
```

60

Part III: Larger Language Features

- Initialization
 - Initializer lists
 - Uniform initialization
 - Prevention of narrowing
- Lambdas
- Rvalue references, move semantics, universal references and perfect forwarding (they're all related...)

61

Problem: Limited Initialization of Aggregates in Old C++

```
int main()
{
    // OK, array initializer
    int vals[] = { 10, 100, 50, 37, -5, 999};

    struct Point { int x; int y; };
    Point p1 = {100,100};    // OK, object initializer

    vector<int> v = { 5, 29, 37};    // ERROR in Old C++!

    const int valsize = sizeof vals / sizeof *vals;

    // range ctor OK
    vector<int> v2(vals, vals + valsize);
}
```

62

Initializer Lists

- C++11's `std::initializer_list` supports generalized initialization of aggregates
- It extends Old C++'s array/object initialization syntax to *any* user-defined type

```
vector<int> v = { 5, 29, 37 };    // Fine in C++11...
vector<int> v2 { 5, 29, 37 };    // Don't even need the =

v2 = { 10, 20, 30, 40, 50 };    // not really just for
                                // "initialization" !

template<typename T>
class vector {                  // A peek inside a typical STL
public:                          // container's implementation...
    vector(std::initializer_list<T>);    // (simplified)
    vector &operator=(std::initializer_list<T>);
    ...
```

63

More Initializer Lists

```
vector<int> foo()
{
    vector<int> v {10, 20, 30};
    v.insert(end(v), { 40, 50, 60 }); // use with algos,

    for (auto x : { 1, 2, 3, 4, 5 })    // with for loops,
        cout << x << " ";
    cout << endl;

    return { 100, 200, 300, 400, 500 }; // most anywhere!
}

int main()
{
    for (auto x : foo())                // note: foo()
        cout << x << " ";            // returns vector
    cout << endl;
}
```

64

Danger, Will Robinson!

- Constructors taking an `initializer_list` are preferred over ones that don't...

```
vector<int> v(10, 20); // Unambiguous: v gets
                      // 10 elements each
                      // equal to 20
```

```
vector<int> v2{10, 20}; // Surprise! v2 gets
                       // two elements: the
                       // values 10 and 20
```

- Advice: going forward, avoid writing constructors that ambiguously overload with others taking an `initializer_list`.

65

Problem: Old Initialization Syntax Can Be Confusing/Ambiguous

```
int main()
{
    int *pi1 = new int(10); // OK, initialized int
    int *pi2 = new int;     // OK, uninitialized
    int *pi3 = new int();   // Now initialized to 0
    int v1(10);             // OK, initialized int
    int v2();               // Oops!

    int foo(bar);           // what IS that?

    int i(5.5);             // legal, unfortunately
    double x = 10e19;
    int j(x);               // even if impossible!
}
```

66

Fix: C++11 Uniform Initialization, Prevention of Narrowing

```
typedef int bar;

int main()
{
    int *pi1 = new int{10}; // initialized int
    int v1{10};             // same
    int *pi2 = new int;     // OK, uninitialized
    int v2{};               // Now it's an object!
    int foo(bar);           // func. declaration
    int foo{bar};           // ERROR with braces
                           // (as it should be)

    double x = 10e19;
    int j{x};               // ERROR: No narrowing
                           // when using {}s

    int i{5.5};             // ERROR, no truncation
}
```

67

Careful When Initializing Aggregates...

- Surprise! Narrowing/truncation of aggregates (arrays and structures) is *always* an error in C++11... even using *legal* Old C++ syntax !!
- This is a breaking change in C++11 (NP, IMO)

```
struct S
{
    int i, j;
};

int main()
{
    S s1 = {5,10};          // OK everywhere
    S s2{5, 10};            // OK in C++11
    S s3{5, 10.5};          // Error everywhere
    S s4 = {5, 10.5};       // ERROR in C++11...
                           // but OK in Old C++!

    int a[] = {1, 2, 3.3};  // ERROR in C++11,
                           // OK in Old C++
}
```

68

More Danger, auto Style

- Another gotcha: when combining initializer lists with `auto`, the type deduced by `auto` is the type of the *initializer list itself*:

```
auto a(10);           // As expected...
cout << "a has type: " <<
    typeid(a).name() << endl;

auto b {10};          // Surprise!
cout << "b has type: " <<
    typeid(b).name() << endl;

// Output:
// a has type: i
// b has type: St16initializer_listIiE
```

69

Problem: Algorithms Not Efficient When Used with Function Pointers

- Inlining rarely applies to function pointers

```
inline bool isPos(int n) { return n > 0; }

int main()
{
    vector<int> v {-5, -19, 3, 10, 15, 20, 100};
    // Calls to isPos probably NOT inlined:
    auto firstPos = find_if(begin(v), end(v), isPos);
    if (firstPos != end(v))
        cout << "First positive value in v is: "
            << *firstPos << endl;

    // Old function object adaptors can eliminate
    firstPos = find_if(begin(v), end(v), // some functions,
        bind2nd(greater<int>(), 0) ); // but they're messy!
}
```

70

Function Objects Improve Performance, But Not Clarity

```
// Have to define a separate class to create function
// objects from:

struct IsPos
{
    bool operator()(int n) { return n > 0; }
};

int main()
{
    vector<int> v {-5, -19, 3, 10, 15, 20, 100};

    auto firstPos =
        find_if(begin(v), end(v), IsPos());
    if (firstPos != end(v))
        cout << "First positive value in v is: "
              << *firstPos << endl;
}
```

71

Lambda Expressions

- A *lambda expression* specifies an anonymous, on-demand function object
- Allows the logic to be truly localized
- Herb Sutter says: "Lambdas make the existing STL algorithms roughly 100x more usable."

```
int main()
{
    vector<int> v {-5, -19, 3, 10, 15, 20, 100};

    auto firstPos = find_if(begin(v), end(v),
        [](int n){return n > 0; });

    if (firstPos != end(v))
        cout << "First positive value in v is: "
              << *firstPos << endl;
}
```

72

Lambdas and Local Variables

- Local variables in scope before the lambda may be *captured* in the lambda's `[]`s
 - The resulting (anon.) function object is called a *closure*

```
int main()
{
    vector<double> v { 1.2, 4.7, 5, 9, 9.4};
    double target = 4.9;
    double epsilon = .3;

    auto endMatches = partition(begin(v), end(v),
                               [target,epsilon] (double val)
                               { return fabs(target - val) < epsilon; });

    cout << "values within epsilon: ";
    for_each(begin(v), endMatches,
             [](double d) { cout << d << ' '; });
    // output: 4.7 5
}
```

73

Different Capture Modes

- Lambdas may capture by reference:

`[&variable1, &variable2]`
- Mix capturing by value and by ref:

`[variable1, &variable2]`
- Specify a default capture mode:

`[=]` (or) `[&]`
- Specify a default, plus special cases:

`[=, &variable1]`

74

Only *Locals* Can Be Captured

- Capturing only applies to *non-static local variables* (including parameters)
- Within a member function, data members cannot be captured directly
 - They may be accessed (and modified) from within the lambda by capturing `this` (even if by value!)
 - Alternatively, they may be copied into a local variable
 - Then the lambda can capture that variable

75

“Avoid Default Capture Modes”

- A lambda may access any global/static data already in scope without any special fanfare
 - Thus, data outside the closure may be modified despite the fact it was not captured
- The potential confusion over what the lambda can modify, along with the potential for dangling references, has moved Scott Meyers to recommend (in his soon-to-be-published *Effective Modern C++*) that programmers “avoid default capture modes.”

76

Lambdas as “Local Functions”

- Defining functions directly within a block is not supported in C++, *but...*

```
int main()
{
    double target = 4.9;
    double epsilon = .3;

    bool withinEpsilonBAD(double val)    // ERROR!
    { return fabs(target - val) < epsilon; };

    auto withinEpsilon = [=](double val)  // OK!
    { return fabs(target - val) < epsilon; };

    cout << (withinEpsilon(5.1) ? "Yes!" : "No!");
}                                           // Output: Yes! 77
```

C++14 Generic Lambdas

```
vector<shared_ptr<string>> vps;

// Example #1:
sort(begin(vps), end(vps), []                // C++11
    (const shared_ptr<string> &p1, const shared_ptr<string> &p2)
    { return *p1 < *p2 } );

sort(begin(vps), end(vps), []                // C++14
    (const auto &p1, const auto &p2) { return *p1 < *p2 });

// Example #2:
auto getsize = []                            // C++11
    (const vector<shared_ptr<string>> &v) { return v.size(); };

auto getsize = []( auto const& c )           // C++14
    { return c.size(); };

// Note: Examples based on Herb's 4/20/13 Trip Report 78
```

Uses for Lambdas

- As we've seen, they're great for use with STL algorithms
 - Predicates for `*_if` algorithms
 - Algos using comparison functions (`sort`, etc.)
- Quick custom deleters for `unique_ptr` and `shared_ptr`
- Easy specification of predicates for condition variables in the threading API
- On-the-fly callback functions

79

Problem: Gratuitous Copying

- In Old C++, objects are (or might be) *copied* when replication is neither needed nor wanted
 - The “extra” copying can sometimes be optimized away (e.g., the RVO), but often is not or cannot

```

class Big { ... };           // expensive to copy

Big makeBig() { return Big(); } // return by value
Big operator+(const Big &, const Big&); // arith. op.

Big bt = makeBig();          // This may cost up to 3
                             // ctors and 2 dtors!

Big x(...), y(...);
Big sum = x + y; // extra copy of ret val from op+ ?

```

80

Old C++ Solutions are Fragile

- The functions *could* be re-written to return:
 - References – but how is memory managed?
 - Raw pointers – prone to leaks, bugs
 - Smart pointers – more syntax and/or overhead
- But if we know the returned object is a *temporary*, we know its data will no longer be needed after “copying” from it
- The solution begins with a new type of reference. But first, some terminology...

81

Ancient Terms, Modern Meanings

- Lvalues
 - Things you can take the address of
 - They may or may not have a name
 - E.g., an expression `*ptr` has no name, but has an address, so it's an lvalue.
- Rvalues
 - Things you can't take the address of
 - Usually they have no name
 - E.g., literal constants, temporaries of various kinds

82

C++11 Rvalue References

- An *rvalue reference* is declared with `&&`
- Binds *only* to (unnamed) temporary objects

```
int fn();                // Note: return val is rvalue
int main()
{
    int i = 10, &ri = i;  // ri is ordinary lvalue ref
    int &&rri = 10;        // OK, rvalue ref to temp
    int &&rri2 = i;        // ERROR, attempt to bind
                        //      lvalue to rvalue ref
    int &&rri3 = i + 10;    // Fine, i + 10 is a temp

    int &ri2 = fn();       // ERROR, attempt to bind
                        //      rvalue to lvalue ref
    const int &ri3 = fn(); // OK, lvalue ref-to-const

    int &&rri4 = fn();      // Fine, ret. val is a temp
}
```

83

Copy vs. Move Operations

- C++ has always had the “copy” operations--the *copy constructor* and *copy assignment operator*:


```
T::T(const T&);           // copy ctor
T &operator=(const T&);   // copy assign.
```
- C++11 adds “move” operations—the *move constructor* and *move assignment operator*:
 - These operations *steal* data from the argument, transfer it to the destination--leaving the argument an “empty husk”
 - This husk can be destroyed / assigned to, and not much else

```
T::T(T &&);               // move ctor
T &operator=(T &&);       // move assignment

// Note: Both would typically be noexcept
```

84

“Big” Class with Move Operations

- So there are now six canonical functions per class (used to be four) that class authors may define

```
class Big {
public:
    Big();                // 1. default ctor
    ~Big();               // 2. destructor
    Big(int x);           // (non-canonical)

    Big(const Big &);      // 3. copy ctor
    Big &operator=(const Big &); // 4. copy assign.
    Big(Big &&);           // 5. move ctor
    Big &operator=(Big &&); // 6. move assign.
private:
    Blob b;               // e.g. some resource-managing type
    double x;             // other data...
};
```

85

Move Operations In Action

```
Big operator+(const Big &, const Big &);
Big munge(const Big &);
Big makeBig() { return Big(); }

int main()
{
    Big x, y;                // Note: below, “created” really
    Big a;                   //      means “not just moved”

    a = makeBig();           // 1 Big created *
    Big b(x + y);            // 1 Big created *
    a = x + y;               // 1 Big created *
    a = munge(x + y);        // 2 Bigs created *
    std::swap(x,y);          // 0 Bigs created!
}

// *: Return value’s contents moved to destination obj
```

86

Move Operations: Not Always Automatic

- Consider the Old C++-style implementation of the `std::swap` function template:

```
template<typename T>
void swap(T &x, T &y)    // lvalue refs
{
    T tmp(x);           // copy ctor
    x = y;               // copy assignment
    y = tmp;             // copy assignment
}
```

- Even when applied to objects (e.g., `Big`) *with move support*, that support won't be used! 87

Forcing Move Operations

- Here's a C++11 version of `std::swap`:

```
template<typename T>
void swap(T &x, T &y)    // still lvalue refs
{
    T tmp(move(x));      // move ctor
    x = move(y);         // move assignment
    y = move(tmp);       // move assignment
}
```

- `move` is a zero-cost function meaning “cast to rvalue”
- Note: this `swap`'s signature is still the same as for old `swap`, but move operations are used *if available* (else falling back on copy operations) 88

Implementing `Big`'s Move Operations

```
class Big {
public:
    ...
    Big(Big &&rhs) :                // move ctor
        b(move(rhs.b)), x(rhs.x) {}

    Big &operator=(Big &&rhs)        // move assignment op.
    {
        b = move(rhs.b);           // Note we NEED the moves, because
        x = rhs.x;                 // rhs itself is an lvalue! (even
        return *this;              // though it has type rvalue ref)
    }

private:
    Blob b;
    double x;
};
```

- `Big`'s move operations simply delegate to `Blob`'s move ops, and assume they do the right thing... 89

`Blob`'s Move Operations

- ...so `Blob`'s move ops must do the “stealing”:

```
class Blob {
public:
    ...
    Blob(Blob &&rhs) {                // move ctor
        raw_ptr = rhs.raw_ptr;       // “steal” pointer
        rhs.raw_ptr = nullptr;       // clear source
    }

    Blob &operator=(Blob &&rhs) {     // move assign. Op
        if (this != &rhs) {
            delete [] raw_ptr;
            raw_ptr = rhs.raw_ptr;    // “steal” pointer
            rhs.raw_ptr = nullptr;    // clear source
        }
        return *this;
    }

private:
    char *raw_ptr;
};
```

90

When && “Doesn’t Mean Rvalue”

- Scott Meyers coined the term *Universal References* for refs--declared using && in a *type deduction* context--that behave as either lvalue or rvalue references:

```

auto &&x = 3.1415;           // x is an rvalue

double pi = 3.14;
auto &&y = pi;               // y is an lvalue

template<typename T>
void f(T &&val);            // Here, val can be
                           // lvalue OR rvalue!

                           // functions instantiated:
f(3.14);                    // f(double &&);
f(x);                      // f(double &&);
f(pi);                     // f(double &);

```

91

Explanation: Reference “Collapsing”

- Refs-to-refs in a universal ref. (*deduction*) context:

- T & & → T&
 - T && & → T&
 - T & && → T&
 - T && && → T&&
- “Lvalue references
are infectious”
-STL

```

template<typename t>
void f(T &&val);            // Here, val can be
                           // lvalue OR rvalue!
double pi = 3.14;

f(3.14);                   // f(double && &&); →
                           //      f(double &&);

f(pi);                     // f(double & &&); →
                           //      f(double &);

```

92

Efficient Sub-object Initialization?

- How many constructors would be required when there are many expensive-to-copy sub-objects?

```
class Big {
public:
    Big(const Blob &b2, const string &str) :    // copy both
        b(b2), s(str) {}

    Big(Blob &&b2, string &&str) :                // move both
        b(move(b2)), s(move(str)) {}

    Big(const Blob &b, string &&str) :           // copy 1st,
        b(b2), s(move(str)) {}                // move 2nd

    Big(Blob &&b, const string &str) :           // move 1st
        b(move(b2)), s(str) {}                // copy 2nd
private:
    Blob b; // what if we added other data members? Arghhh!
    string s;
};
```

93

Perfect Forwarding

- We'd prefer for each parameter to be copied or moved *as per its original lvalue-ness or rvalue-ness*

```
class Big {
public:
    template<typename T1, typename T2>
    Big(T1 &&b2, T2 &&str) :    // Universal refs
        b(std::forward<T1>(b2), // std::forward preserves the
            s(std::forward<T2>(str)) // lvalue-ness or rvalue-ness
        {}                    // (and const-ness) of its arg

private:
    Blob b;
    string s;
};
```

94

When Move-Enable a Type?

- In the general case, move operations should be added only when *moving can be implemented faster than copying*
- Most C++11 library components are move-enabled
 - Some (e.g. `unique_ptr`, covered later) are *move-only*--they don't support conventional copy operations.
 - Internally, the implementations of many components, e.g. containers, employ moves whenever possible (rather than copying)

95

“The Rule of 5”

- The Old C++ “Rule of 3” now becomes the “Rule of 5”:
- Good C++11 style dictates that if you declare any copy operation, move operation or destructor (even if only with `=default` or `=delete`), then you should declare all 5
- If you declare *any* of the 5, *no* move operations will be generated implicitly
 - The *copy* operations are still generated if needed
 - Note, however: this behavior is *deprecated in C++11!*

96

C++14 Generalized Lambda Capture

- C++11 lambdas can only capture variables by value or by reference
 - So those lambdas can't capture move-only types
 - C++14's can be initialized with *arbitrary expressions*
 - They can have their own names, but even if not, the captured values are distinct from the originals

```
// Example from C++14 wikipedia page:  
auto ptr = std::make_unique<int>(10);  
auto lambda =  
    [ptr = std::move(ptr)] {return *ptr;};
```

97

Epilogue: Is C++ Too Complicated?

- This is an oft-heard complaint
- But how does one measure “complexity”?
 - Pages in the *language* specification?
 - C++98 Language: 309 pp.
 - C++11 Language: 424 pp.
 - How does that compare to other languages?
 - Java SE 7 Edition: 606 pp.
 - C# ECMA Standard 4th Ed. June '06: 531 pp.

98

Complexity vs. “Complications”

- I'd like to suggest some different metrics for what makes a language “complicated to use”:
 - To what extent does good code rely on (undocumented) idioms?
 - How many “gotchas” show up while debugging?
 - Can you “say it in code” or do you have to explain what you’re doing in comments?
 - Does attaining high performance require jumping through hoops?
- By these criteria, I believe C++11/14 “measures up” pretty well!

99

C++ Resources

- For live links to resources listed here and more, please visit my “links” page at BD Software:
www.bdsoft.com/links.html
- The C++ Standards Committee:
www.open-std.org/jtc1/sc22/wg21
(Draft C++ Standard available for free download)
- ISO C++ Site (spearheaded by Herb Sutter and the Standard C++ Foundation):
isocpp.org

100

Overviews of C++11/14

- Bjarne Stroustrup's C++11 FAQ:
<http://www.stroustrup.com/C++11FAQ.html>
- Wikipedia C++11 page:
en.wikipedia.org/wiki/C++11
- Elements of Modern C++ Style (Herb Sutter):
herbsutter.com/elements-of-modern-c-style/
- Scott Meyers' *Overview of the New C++ (C++11)*
[http://www.artima.com/shop/
overview_of_the_new_cpp](http://www.artima.com/shop/overview_of_the_new_cpp)

101

C++14 Features

- *C++14 and early thoughts about C++17* (Bjarne Stroustrup):
[https://parasol.tamu.edu/people/bs/
622-GP/C++14TAMU.pdf](https://parasol.tamu.edu/people/bs/622-GP/C++14TAMU.pdf)
- *A Look at C++14: Papers* (Meeting C++)
[http://www.meetingcpp.com/index.php/br/items/
a-look-at-cpp14-papers-part-1.html](http://www.meetingcpp.com/index.php/br/items/a-look-at-cpp14-papers-part-1.html)
- C++14 Wiki
<http://en.wikipedia.org/wiki/C++14>

102

On Specific New C++ Features

- *Rvalue References and Perfect Forwarding Explained* (Thomas Becker):
http://thbecker.net/articles/rvalue_references/section_01.html
- *Universal References in C++* (Scott Meyers)
 - Article, with link to great video from C&B '12:
<http://isocpp.org/blog/2012/11/universal-references-in-c11-scott-meyers>
- *Lambdas, Lambdas Everywhere* (Herb Sutter)
 - These are the slides (there are videos out there too):
<http://tinyurl.com/lambda-lambdas>

103

On Specific New C++ Features: Using the Standard Smart Pointers

- *Guru of the Week #89: Smart Pointers* (Herb Sutter):
<http://herbsutter.com/2013/05/29/gotw-89-solution-smart-pointers>
- *Guru of the Week #91: Smart Pointer Parameters* (Herb Sutter):
<http://herbsutter.com/2013/06/05/gotw-91-solution-smart-pointer-parameters>

104

Multimedia Presentations

- Herb Sutter
 - *Why C++?* (Herb's amazing keynote from C++ *and Beyond 2011*, a few days before C++11's ratification):
channel9.msdn.com/posts/C-and-Beyond-2011-Herb-Sutter-why-C
 - *Writing modern C++ code: how C++ has evolved over the years*:
channel9.msdn.com/Events/BUILD/BUILD2011/TOOL-835T
- Going Native 201x (@ µSoft) Talks
 - Bjarne, Herb, Andre, "STL", many others:
<http://channel9.msdn.com/Events/GoingNative>

105

Concurrency Resources

- Tutorials
 - Book: *C++ Concurrency in Action* (Williams)
 - Tutorial article series by Williams:
Multithreading in C++0x (parts 1-8)
 - *C++11 Concurrency Series* (9 videos, Milewski)
- `just::thread` Library Reference Guide
 - www.stdthread.co.uk/doc

106

Where to Get Compilers / Libraries

- Twilight Dragon Media (TDM) gcc compiler for Windows
tdm-gcc.tdragon.net/start
- Visual C++ Express compiler
<http://www.microsoft.com/visualstudio/eng/downloads>
- Boost libraries
www.boost.org
- Just Software Solutions (just::thread library)
www.stdthread.co.uk
- If running under Cygwin, a Wiki on building the latest gcc distro under that environment:
http://cygwin.wikia.com/wiki/How_to_install_a_newer_version_of_GCC

107

Testing Code on PCs

- Here are some good bets for trying out new C++ features on a Windows platform:
 - g++ 4.9.1 distro at *nuwen.net* (courtesy of Microsoft's own "STL"). Note: no threading support, but lots of C++14 language support:
<http://nuwen.net/mingw.html>
 - clang 3.6 via the online compiler at *Wandbox* (full language/lib support):
<http://melpon.org/wandbox>

108

"There are only two kinds of languages: the ones people complain about and the ones nobody uses."

-Bjarne Stroustrup

Thanks for attending!

Leor Zolman
leor@bdsoft.com

For many of the links cited here, and more, please
visit: www.bdsoft.com/links.html