
C++11 in the Wild

Techniques from a real codebase

What we'll cover

- The **Auto()** macro

(slides 3–34)

run arbitrary cleanup code at the end of a scope

- **make_iterable()** and **iterator_range**

(slides 35–49)

turn a pair of iterators into a container

Time permitting: • **std::spaceship**

(slides 50–81)

compare any two objects strcmp-wise

Chapter 1.

The Auto() macro

A better OnScopeExit()

What we start with

```
void Mutate(State *state)
{
    state->DisableLogging();
    state->AttemptOperation();
    state->AttemptDifferentOperation();
    state->EnableLogging();
    return;
}
```

Oops, forgot all the error handling

```
bool Mutate(State *state)
{
    state->DisableLogging();
    if (!state->AttemptOperation()) return false;
    if (!state->AttemptDifferentOperation()) return false;
    state->EnableLogging();
    return true;
}
```

(Or, use exceptions for control flow if you want.
You'll have the same problem.)

What we want to write

```
#include "auto.h"
```

```
bool Mutate(State *state)
{
    state->DisableLogging();
    Auto(state->EnableLogging());

    if (!state->AttemptOperation()) return false;
    if (!state->AttemptDifferentOperation()) return false;
    return true;
}
```

#include "auto.h"

Credits: Marko Tintor, Alex Skidanov, Arthur O'Dwyer

```
#pragma once

template <class Lambda> class AtScopeExit {
    Lambda& m_lambda;
public:
    AtScopeExit(Lambda& action) : m_lambda(action) {}
    ~AtScopeExit() { m_lambda(); }
};

#define TOKEN_PASTE(x, y) x ## y
#define TOKEN_PASTE(x, y) TOKEN_PASTE(x, y)

#define Auto_INTERNAL1(lname, aname, ...) \
    auto lname = [&]() { __VA_ARGS__; }; \
    AtScopeExit<decltype(lname)> aname(lname);

#define Auto_INTERNAL2(ctr, ...) \
    Auto_INTERNAL1(TOKEN_PASTE(Auto_func_, ctr), \
        TOKEN_PASTE(Auto_instance_, ctr), __VA_ARGS__)

#define Auto(...) Auto_INTERNAL2(__COUNTER__, __VA_ARGS__)
```

Choose Your Own Digression

- Variadic macros and `__VA_ARGS__` (C++11)
 - Token pasting and `##`
 - Templates
 - Lambdas (C++11)
 - `__COUNTER__` (non-standard)
 - `#pragma once` (non-standard)
 - Style point: Aren't macros evil or something?
 - Style point: Why lambdas instead of `std::function`?
-

__COUNTER__

It gives a new integer value every time it's expanded.

It's non-standard,
but every compiler in the world supports it.

I almost said ***almost*** every compiler,
but I can't name any compilers that don't support it.

We would avoid it if there were any standard way to get its functionality.

__LINE__ kinda works, I guess...

#pragma once

#pragma once is the clearest, most efficient way to make a file idempotent.

It's non-standard,
but every compiler in the world supports it.

#ifndef, #define, and #endif have their own uses,
but you don't need them (and therefore shouldn't
use them) to make a file idempotent.

Idempotence

A function $f: D \rightarrow D$ is idempotent if

$$f(fx) = fx \text{ for all } x \text{ in } D.$$

I.e., repeated applications have the same effect as one.

(FOLDOC)

Why not std::function?

```
template <class Lambda> class AtScopeExit {  
    Lambda& m_lambda;  
public:  
    AtScopeExit(Lambda& action) : m_lambda(action) {}  
    ~AtScopeExit() { m_lambda(); }  
};
```

```
AtScopeExit<decltype(lname)> aname(lname);
```

green text: what we wrote

```
class AtScopeExit {  
    std::function<void(void)> m_lambda;  
public:  
    template<class Lambda> AtScopeExit(Lambda& action) : m_lambda(action) {}  
    ~AtScopeExit() { m_lambda(); }  
};
```

```
AtScopeExit aname(lname);
```

red text: what we consciously
chose not to write

Why not `std::function`?

- We don't want to pull in all of `<functional>`.
 - “auto.h” is included by generated code and must be lightweight.
 - `std::function` uses type erasure, which uses heap allocation.
 - More on this later.
 - Empirically, we get better code this way.
 - Assembly listings on next page.
-

Let's see some codegen!

```
#include <stdio.h>
#include "auto.h"

extern void foo();

int main() {
    if (true) {
        Auto(puts("two"));
        puts("one");           // compiler knows this doesn't throw
    }
    if (true) {
        Auto(puts("three"));
        foo();                 // might throw an exception
    }
}
```

Let's see some codegen!

Clang 3.4 -O2 gives
perfect code

```
#include <stdio.h>
#include "auto.h"
```

```
extern void foo();
```

```
int main() {
    if (true) {
        Auto(puts("two"));
        puts("one");
    }
    if (true) {
        Auto(puts("three"));
        foo();
    }
}
```

```
_main:
    pushq %rbp
    movq %rsp, %rbp
    pushq %rbx
    pushq %rax
    leaq L_.str(%rip), %rdi ## "one"
    callq _puts
    leaq L_.str2(%rip), %rdi ## "two"
    callq _puts
    callq __Z3foov
    ## reached iff foo doesn't throw any exception
    leaq L_.str1(%rip), %rdi ## "three"
    callq _puts
    xorl %eax, %eax
    addq $8, %rsp
    popq %rbx
    popq %rbp
    ret
LBB0_2: ## reached iff foo throws an exception
    movq %rax, %rbx
    leaq L_.str1(%rip), %rax ## "three"
    movq %rax, %rdi
    callq _puts
    movq %rbx, %rdi
    callq __Unwind_Resume
```

To remove this stack frame, use -O3.

Let's see some codegen!

GCC 4.8 -O2 gives
perfect code
(but only if you give it a hint)

```
extern void puts(const char*)
    noexcept(true); Clang is smarter than GCC
about the standard library's
noexcept guarantees.

extern void foo();

int main() {
    if (true) {
        Auto(puts("two"));
        puts("one");
    }
    if (true) {
        Auto(puts("three"));
        foo();
    }
}
```

```
main:
    pushq    %rbx
    movl     $.LC0, %edi    ## "one"
    call     _Z4putsPKc
    movl     $.LC1, %edi    ## "two"
    call     _Z4putsPKc
    call     _Z3foov
    ## reached iff foo doesn't throw any exception
    movl     $.LC2, %edi    ## "three"
    call     _Z4putsPKc
    xorl     %eax, %eax
    popq     %rbx
    ret
.L3:  ## reached iff foo throws an exception
    movq     %rax, %rbx
    movl     $.LC2, %edi    ## "three"
    call     _Z4putsPKc
    movq     %rbx, %rdi
    call     _Unwind_Resume
```


Let's see some codegen!

The std::function version
is objectively terrible.

```

                _main:
#include <stdio.h>    pushq %rbp
#include "auto.h"     movq %rsp, %rbp
                     ## BB#1:
                     pushq %r14    leaq -80(%rbp), %rdi
                     pushq %rbx    callq __ZN15AtScopeExitD2Ev
                     subq $64, %rsp  movq (%r14), %rax
extern void foo();    movq ___stack_chk_guard@GOTPCREL(%rip), %r14  cmpq -24(%rbp), %rax
                     jne LBB0_4
                     movq (%r14), %rax  ## BB#2:
                     movq %rax, -24(%rbp)  xorl %eax, %eax
                     leaq -80(%rbp), %rbx  addq $64, %rsp
                     movq %rbx, -48(%rbp)  popq %rbx
                     leaq __ZTVNSt3__110__function6__f  popq %r14
                     movq %rax, -80(%rbp)  popq %rbp
                     leaq L_.str(%rip), %rdi  ret
                     callq _puts  LBB0_4:
                     movq %rbx, %rdi  callq ___stack_chk_fail
                     callq __ZN15AtScopeExitD2Ev  LBB0_3:
                     movq %rbx, -48(%rbp)  movq %rax, %rbx
                     leaq __ZTVNSt3__110__function6__f  leaq -80(%rbp), %rax
                     movq %rax, -80(%rbp)  movq %rax, %rdi
                     callq __Z3foov  callq __ZN15AtScopeExitD2Ev
                                     movq %rbx, %rdi
                                     callq __Unwind_Resume

int main() {
    if (true) {
        Auto(puts("two"
        puts("one");
    }
    if (true) {
        Auto(puts("thre
        foo();
    }
}
```

Let's see some codegen!

The std::function version
is objectively terrible.

```

                _main:
#include <stdio.h>    pushq %rbp
#include "auto.h"     movq %rsp, %rbp
                     ## BB#1:
                     pushq %r14    leaq -80(%rbp), %rdi
                     pushq %rbx    callq __ZN15AtScopeExitD2Ev
                     subq $64, %rsp movq (%r14), %rax
extern void foo();    movq ___stack_chk_guard@GOTPCREL(%rip), %r14 cmpq -24(%rbp), %rax
                     jne LBB0_4
                     movq (%r14), %rax
                     ## BB#2:
                     movq %rax, -24(%rbp)    xorl %eax, %eax
                     leaq -80(%rbp), %rbx    addq $64, %rsp
                     movq %rbx, -48(%rbp)    popq %rbx
                     leaq __ZTVNSt3__110__function6__f popq %r14
                     movq %rax, -80(%rbp)    popq %rbp
                     leaq L_.str(%rip), %rdi    ret
                     callq _puts
                     LBB0_4:
                     movq %rbx, %rdi    callq ___stack_chk_fail
                     callq __ZN15AtScopeExitD2Ev
                     LBB0_3:
                     movq %rbx, -48(%rbp)    movq %rax, %rbx
                     leaq __ZTVNSt3__110__function6__f leaq -80(%rbp), %rax
                     movq %rax, -80(%rbp)    movq %rax, %rdi
                     callq __Z3foov    callq __ZN15AtScopeExitD2Ev
                                     movq %rbx, %rdi
                                     callq __Unwind_Resume

int main() {
    if (true) {
        Auto(puts("two"
        puts("one");
    }
    if (true) {
        Auto(puts("thre
        foo();
    }
}

(700 lines of std::function code omitted)
```

**So how is `std::function` implemented,
to get such bad performance?**

Type erasure in a nutshell

To capture any type:

(1) Make a Container that can hold any type.

I.e., make a template class.

```
template<typename T> class Container
{
    T captured_object;
}
```

Type erasure in a nutshell

To capture any type:

(2) Make a `TypeErasedObject` that can hold `Container<T>` for any `T`.

Via polymorphism (inheritance and virtual dispatch).

```
template <typename T> class Container : ContainerBase;
class TypeErasedObject {
    ContainerBase *container;
    TypeErasedObject(X x) { container = new Container<X>(x); }
};
```

#include "function.h"

```
#pragma once
#include <utility>

struct ContainerBase {
    virtual void perform() = 0;
    virtual ~ContainerBase() = default;
};

template <class Lambda> struct Container : ContainerBase {
    Lambda m_lambda;
    Container(Lambda&& lambda) : m_lambda(std::move(lambda)) {}
    virtual void perform() { m_lambda(); }
};

class function { // equivalent to std::function<void(void)>
    ContainerBase *m_ctr;
public:
    template<class Lambda> function(Lambda lambda)
        : m_ctr(new Container<Lambda>(std::move(lambda))) {}
    void operator() () { m_ctr->perform(); }
    ~function() { delete m_ctr; }
};
```

#include "function.h"

```
#pragma once
```

```
#include <utility>
```

std::move has a compile-time cost, as it relies on std::remove_reference

```
struct ContainerBase {  
    virtual void perform() = 0;  
    virtual ~ContainerBase() = default;  
};
```

virtual dispatch has a runtime cost

```
template <class Lambda> struct Container : ContainerBase {  
    Lambda m_lambda;  
    Container(Lambda&& lambda) : m_lambda(std::move(lambda)) {}  
    virtual void perform() { m_lambda(); }  
};
```

we cannot avoid move-constructing a Lambda here; this move-constructs all its captures (but in our case this is cheap, because we captured them by reference)

```
class function { // equivalent to std::function<void(void)>  
    ContainerBase *m_ctr;  
public:  
    template<class Lambda> function(Lambda lambda)  
        : m_ctr(new Container<Lambda>(std::move(lambda))) {}  
    void operator()() { m_ctr->perform(); }  
    ~function() { delete m_ctr; }  
};
```

memory allocation has a huge runtime cost, although we may avoid it if sizeof (Lambda) is small (via a kind of “small string optimization”)

Alternative syntaxes

- Alexandrescu & Marginean's ScopeGuard
 - Boost.ScopeExit
 - Google scope-exit
-

Alexandrescu & Marginean

Generic: Change the Way You Write Exception-Safe Code — Forever

Andrei Alexandrescu and Petru Marginean, December 2000

<http://www.drdobbs.com/cpp/generic-change-the-way-you-write-excepti/184403758>

```
ScopeGuard guard = MakeObjGuard(state, &State::EnableLogging);
```

```
ON_BLOCK_EXIT(state, &State::EnableLogging);
```

Alexandrescu & Marginean

Generic: Change the Way You Write Exception-Safe Code — Forever

Andrei Alexandrescu and Petru Marginean, December 2000

<http://www.drdobbs.com/cpp/generic-change-the-way-you-write-excepti/184403758>

```
ScopeGuard guard = MakeObjGuard(state, &State::EnableLogging);
```

```
ON_BLOCK_EXIT(state, &State::EnableLogging);
```

Can't run arbitrary code unless it's wrapped in a function

Can't write your cleanup code in-line with your other code

Cleanup code can't refer to local variables

Boost.ScopeExit

Plain vanilla Boost:

```
BOOST_SCOPE_EXIT(&state) {  
    state->EnableLogging();  
} BOOST_SCOPE_EXIT_END
```

Or, if you have C++11, Boost provides:

```
BOOST_SCOPE_EXIT_ALL(&) { state->EnableLogging(); };
```

Or, a C++11 alternative suggested in the Annex:

```
scope_exit on_exit42([&]{ state->EnableLogging(); });
```

Boost.ScopeExit

Plain vanilla Boost:

```
BOOST_SCOPE_EXIT(&state) {  
    state->EnableLogging();  
} BOOST_SCOPE_EXIT_END
```

Or, if you have C++11, Boost provides:

```
BOOST_SCOPE_EXIT_ALL(&) { state->EnableLogging(); };
```

Or, a C++11 alternative suggested in the Annex:

```
scope_exit on_exit42([&]{ state->EnableLogging(); });
```

Very similar to Auto(), but so much boilerplate!

scope_exit requires coming up with unique names (not friendly to code-generation)

Google scope-exit

```
ON_SCOPE_EXIT((state), state->EnableLogging());
```

An example from their documentation:

```
template<typename T>
void f(T& t)
{
    int i, x;

    ON_SCOPE_EXIT((i) SCOPE_EXIT_TEMPLATE_VAR(t) (x),
        /* Do something with i, t, and x */
    );
}
```

Google scope-exit

```
ON_SCOPE_EXIT((state), state->EnableLogging());
```

An example from their documentation:

```
template<typename T>
void f(T& t)
{
    int i, x;

    ON_SCOPE_EXIT((i) SCOPE_EXIT_TEMPLATE_VAR(t) (x),
        /* Do something with i, t, and x */
    );
}
```

Must explicitly name all your captures (unfriendly to code-generation)

Weird corner cases with templates and the “this” pointer

One more time

```
#pragma once
```

```
template <class Lambda> class AtScopeExit {  
    Lambda& m_lambda;  
public:  
    AtScopeExit(Lambda& action) : m_lambda(action) {}  
    ~AtScopeExit() { m_lambda(); }  
};
```

```
#define TOKEN_PASTE(x, y) x ## y  
#define TOKEN_PASTE(x, y) TOKEN_PASTE(x, y)
```

```
#define Auto_INTERNAL1(lname, aname, ...) \  
    auto lname = [&]() { __VA_ARGS__; }; \  
    AtScopeExit<decltype(lname)> aname(lname);
```

```
#define Auto_INTERNAL2(ctr, ...) \  
    Auto_INTERNAL1(TOKEN_PASTE(Auto_func_, ctr), \  
        TOKEN_PASTE(Auto_instance_, ctr), __VA_ARGS__)
```

```
#define Auto(...) Auto_INTERNAL2(__COUNTER__, __VA_ARGS__)
```

One odd application

```
CodePrinter& code = context.codeprinter;

code.Printf("void MergeWith(OtherRowElement* other, const TableColumns_%s*
/*dummy*/, int threadId)\n", ti[i].tableAlias);
code.Scope();
code.Printf("if (other->%s == nullptr)\n", ti[i].tableResultName);
code.Scope();
code.Printf("%s = nullptr;\n", ti[i].tableResultName);
code.Unscope();
code.Printf("else\n");
code.Scope();
CodeGenElseBlock(context, ti, i);
code.Unscope();
code.Unscope(); // end of function body
```

One odd application

```
#define AutoScope(code) code.Scope(); Auto(code.Unscope());
```

```
code.Printf("void MergeWith(OtherRowElement* other, const TableColumns_%s*\n", ti[i].tableAlias);\n/*dummy*/, int threadId)\n{\n    AutoScope(code);\n    code.Printf("if (other->%s == nullptr)\n", ti[i].tableResultName);\n    {\n        AutoScope(code);\n        code.Printf("%s = nullptr;\n", ti[i].tableResultName);\n    }\n    code.Printf("else\n");\n    {\n        AutoScope(code);\n        CodeGenElseBlock(context, ti, i);\n    }\n}
```

Any questions?

```
#pragma once
```

```
template <class Lambda> class AtScopeExit {  
    Lambda& m_lambda;  
public:  
    AtScopeExit(Lambda& action) : m_lambda(action) {}  
    ~AtScopeExit() { m_lambda(); }  
};
```

```
#define TOKEN_PASTE(x, y) x ## y  
#define TOKEN_PASTE(x, y) TOKEN_PASTE(x, y)
```

```
#define Auto_INTERNAL1(lname, aname, ...) \  
    auto lname = [&]() { __VA_ARGS__; }; \  
    AtScopeExit<decltype(lname)> aname(lname);
```

```
#define Auto_INTERNAL2(ctr, ...) \  
    Auto_INTERNAL1(TOKEN_PASTE(Auto_func_, ctr), \  
        TOKEN_PASTE(Auto_instance_, ctr), __VA_ARGS__)
```

```
#define Auto(...) Auto_INTERNAL2(__COUNTER__, __VA_ARGS__)
```

Chapter 2.

`make_iterable`

Inside-out containers

What we start with

```
class MDTable
{
    MDColumn *m_columns;
    MDKey *m_keys;
    int m_columnCount;
    int m_keyCount;
public:
    MDColumn* GetColumns() const { return m_columns; }
    int GetNumColumns() const { return m_columnCount; }
    MDIndex* GetKeys() const { return m_keys; }
    int GetNumKeys() const { return m_keyCount; }
};
```

Why not just use std::vector?

```
class MDTable
{
    MDKey *m_keys;    // both normal and foreign
    int m_keyCount;
    int m_firstForeignKey;

public:
    MDIndex* GetNormalKeys() const { return m_keys; }
    int GetNumNormalKeys() const { return m_firstForeignKey; }
    MDIndex* GetForeignKeys() const { return m_keys + ...; }
    int GetNumForeignKeys() const { return m_keyCount - ...; }
};
```

Why not just use std::vector?

```
class MDTable
{
    MDKey *m_keys;    // both normal and foreign keys
    int m_keyCount;
    int m_firstForeignKey;

public:
    MDIndex* GetNormalKeys() const { return m_keys; }
    int GetNumNormalKeys() const { return m_firstForeignKey; }
    MDIndex* GetForeignKeys() const { return m_keys + ...; }
    int GetNumForeignKeys() const { return m_keyCount - ...; }
};
```

Weird design choices.

Time efficiency.

Space efficiency.

Using GetColumns() is cumbersome

```
void TransformTable(MDTable *tab)
{
    for (int i=0; i < tab->GetNumColumns(); ++i)
    {
        MDColumn& col = tab->GetColumns()[i];
        ... col ...
    }
    for (int i=0; i < tab->GetNumKeys(); ++i)
    {
        MDKey& key = tab->GetKeys()[i];
        ... key ...
    }
}
```

What we'd like to write

```
void TransformTable(MDTable *tab)
{
    for (MDColumn& col : Columns(tab))
    {
        ... col ...
    }
    for (MDKey& key : Keys(tab))
    {
        ... key ...
    }
}
```

Our Columns() and Keys() functions

```
#include "iterable.h"
```

```
static inline iterable<MDColumn*> Columns(MDTable* tab)
{
    MDColumn* cols = tab->GetColumns();
    return make_iterable(cols, cols + tab->GetNumColumns());
}
```

```
static inline iterable<MDKey*> Keys(MDTable* tab)
{
    MDKey* keys = tab->GetKeys();
    return make_iterable(keys, keys + tab->GetNumKeys());
}
```

Our Columns() and Keys() functions

```
#include "iterable.h"
```

```
static inline iterable<MDColumn*> Columns(MDTable* tab)
{
    MDColumn* cols = tab->GetColumns();
    return make_iterable(cols, cols + tab->GetNumColumns());
}
```

```
static inline iterable<MDKey*> Keys(MDTable* tab)
{
    MDKey* keys = tab->GetKeys();
    return make_iterable(keys, keys + tab->GetNumKeys());
}
```

#include "iterable.h"

```
template<class It>
class iterable
{
    It m_first, m_last;
public:
    iterable() = default;
    iterable(It first, It last) :
        m_first(first), m_last(last) {}
    It begin() const { return m_first; }
    It end() const { return m_last; }
};

template<class It>
inline iterable<It> make_iterable(It a, It b)
{
    return iterable<It>(a, b);
}
```

#include "iterable.h"

```
template<class It>
class iterable
{
    It m_first, m_last;
public:
    iterable() = default;
    iterable(It first, It last) :
        m_first(first), m_last(last) {}
    It begin() const { return m_first; }
    It end() const { return m_last; }
};
```

Marshall Clow called it `iterator_pair`

<http://cplusplusmusings.wordpress.com/2013/04/14/range-based-for-loops-and-pairs-of-iterators/>

Google calls it `std::range`

<http://cxx1y-range.googlecode.com/git/paper.html>

Boost calls it `iterator_range`

http://www.boost.org/doc/libs/1_53_0/libs/range/doc/html/range/reference/utilities/iterator_range.html

Alisdair Meredith (N2977) called it `std::range`

<http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2009/n2977.pdf>

```
template<class It>
inline iterable<It> make_iterable(It a, It b)
{
    return iterable<It>(a, b);
}
```

It's basically just a pair of iterators...

Original Frankfurt C++11 proposal called it `std::pair`

Original Frankfurt C++11 proposal called it `std::pair`

but this is a bad idea

Original Frankfurt C++11 proposal called it `std::pair`

but this is a bad idea

because there are standard algorithms
that deal in pairs of iterators
that are **not ranges**

Abstract

What are not needed

1. *Journal of the American Medical Association*, 1997; 277: 1039-1043.

```
template<class ForwardIt>
```

```
template<class InputT, class OutputT1, class OutputT2, class UnaryPredicate>
```


Inside-out containers

```
template<class It>
class iterable
{
    It m_first, m_last;
public:
    iterable() = default;
    iterable(It first, It last) :
        m_first(first), m_last(last) {}
    It begin() const { return m_first; }
    It end() const { return m_last; }
};
```

```
template<class It>
inline iterable<It> make_iterable(It a, It b)
{
    return iterable<It>(a, b);
}
```

Make a “container view” of an object on the fly

One object can have multiple iterable parts, without exposing implementation details

You can iterate over a subrange as easily as you iterate over the entire container

Still no word on “ranges” in C++1z (there is a working group)

Chapter 3.

`std::spaceship`

(which doesn't exist)

Motivating use case (LLVM)

```
/// array_pod_sort - This sorts an array with the specified start and end extent.  
/// This is just like std::sort, except that it calls qsort instead of      using an inlined template.  
/// qsort is slightly slower than std::sort, but most sorts are not performance critical in LLVM  
/// and std::sort has to be template instantiated for each type, leading to significant measured  
/// code bloat.  This function should generally be used instead of std::sort where possible.  
///  
/// This function assumes that you have simple POD-like types that can be compared with operator<  
/// and can be moved with memcpy.  If this isn't true, you should use std::sort.  
///
```

```
template <class IteratorTy>  
inline void array_pod_sort(  
    IteratorTy Start,  
    IteratorTy End,  
    int (*Compare) (  
        const typename std::iterator_traits<IteratorTy>::value_type *,  
        const typename std::iterator_traits<IteratorTy>::value_type *) ) {  
    // Don't dereference start iterator of empty sequence.  
    if (Start == End) return;  
    qsort(&*Start, End - Start, sizeof(*Start),  
        reinterpret_cast<int (*) (const void *, const void *)>(Compare));  
}
```

with some of the cruft removed

```
/// array_pod_sort - This sorts an array with the specified start and end extent.
/// This is just like std::sort, except that it calls qsort instead of      using an inlined template.
/// qsort is slightly slower than std::sort, but most sorts are not performance critical in LLVM
/// and std::sort has to be template instantiated for each type, leading to significant measured
/// code bloat.  This function should generally be used instead of std::sort where possible.
///
/// This function assumes that you have simple POD-like types that can be compared with operator<
/// and can be moved with memcpy.  If this isn't true, you should use std::sort.
///

template <typename T>
void array_pod_sort(T *start, T *end, int (*compare)(const T *, const T *))
{
    // Don't dereference start iterator of empty sequence.
    if (start == end) return;
    std::qsort(/* base      */ start,
               /* nelem    */ end - start,
               /* width    */ sizeof *start,
               /* compar   */ reinterpret_cast<int (*) (const void *, const void *)>(compare));
}
```

Some example comparators

```
static int SrcCmp(const std::pair<const CFGBlock *, const Stmt *> *p1,
                  const std::pair<const CFGBlock *, const Stmt *> *p2) {
    if (p1->second->getLocStart() < p2->second->getLocStart())
        return -1;
    if (p2->second->getLocStart() < p1->second->getLocStart())
        return 1;
    return 0;
}
```

```
static int compareEntry(const Table::MapEntryTy *const *LHS,
                       const Table::MapEntryTy *const *RHS) {
    return (*LHS)->getKey().compare((*RHS)->getKey());
}
```

```
static int CompareCXXCtorInitializers(CXXCtorInitializer *const *X,
                                       CXXCtorInitializer *const *Y) {
    return (*X)->getSourceOrder() - (*Y)->getSourceOrder();
}
```

<https://www.mail-archive.com/cfe-commits@cs.uiuc.edu/msg92289.html>
Clang's r203293 "[C++11] Revert uses of lambdas with array_pod_sort."

This is a pretty common idiom

Particularly in C.

`strcmp`

`(strcoll, strcasecmp...)`

`qsort`

`bsearch`

This is a pretty common idiom

But also (occasionally) in C++!

The diagram illustrates the relationship between C++ standard library functions and their headers. It lists four functions on the left and their corresponding headers on the right, with arrows pointing from the headers to the functions:

- `std::string::compare` is associated with `<string>`.
- `std::char_traits<T>::compare` is also associated with `<string>`.
- `std::collate<T>::compare` is associated with `<locale>`.
- `std::sub_match<T>::compare` is associated with `<regex>`.

```
graph LR; S1[std::string::compare] --> H1[<string>]; S2[std::char_traits<T>::compare] --> H1; S3[std::collate<T>::compare] --> H2[<locale>]; S4[std::sub_match<T>::compare] --> H3[<regex>];
```

But look at the variety of these comparators!

```
if (a < b) return -1;           // the Java programmer's approach
if (a > b) return 1;
return 0;

return a.compare(b);           // otherwise known as "delegating the task to someone else"

return a - b;                   // short and sweet, but can lead to bugs

return (a < b) ? -1 : (a > b);   // my personal favorite

return (a < b) ? -1 : (b < a) ? 1 : 0; // the minimalist approach: uses only operator<

if (a != b) return (a < b) ? -1 : 1; // the extensible approach
return 0;
```


Wouldn't it be nice if there were a
simple, unified way
to write comparators like these?

This is a solved problem

In Perl, Ruby, Groovy...

This is a solved problem

In Perl, Ruby, Groovy...

The spaceship operator



This is a solved problem

In Perl, Ruby, Groovy...

The spaceship operator

a <=> b

means

(a < b) ? -1 :

(a > b) ? +1 : 0

This won't be about operator overloading

We don't care how the operation is spelled
(and spelling it \leq seems out of the question)

We can just spell it **std::spaceship**.
That's fine.

First attempt: LLVM to the rescue!

```
/// array_pod_sort_comparator - This is helper function for array_pod_sort,  
/// which just uses operator< on T.  
template<typename T>  
inline int array_pod_sort_comparator(const void *P1, const void *P2) {  
    if (*reinterpret_cast<const T*>(P1) < *reinterpret_cast<const T*>(P2))  
        return -1;  
    if (*reinterpret_cast<const T*>(P2) < *reinterpret_cast<const T*>(P1))  
        return 1;  
    return 0;  
}
```

First attempt: LLVM to the rescue!

```
/// array_pod_sort_comparator - This is helper function for array_pod_sort,  
/// which just uses operator< on T.  
template<typename T>  
inline int array_pod_sort_comparator(const void *P1, const void *P2) {  
    if (*reinterpret_cast<const T*>(P1) < *reinterpret_cast<const T*>(P2))  
        return -1;  
    if (*reinterpret_cast<const T*>(P2) < *reinterpret_cast<const T*>(P1))  
        return 1;  
    return 0;  
}
```

The problem is inefficiency.

First attempt: LLVM to the rescue!

```
/// array_pod_sort_comparator - This is helper function for array_pod_sort,  
/// which just uses operator< on T.  
template<typename T>  
inline int array_pod_sort_comparator(const void *P1, const void *P2) {  
    if (*reinterpret_cast<const T*>(P1) < *reinterpret_cast<const T*>(P2))  
        return -1;  
    if (*reinterpret_cast<const T*>(P2) < *reinterpret_cast<const T*>(P1))  
        return 1;  
    return 0;  
}
```

The problem is inefficiency.

Two calls to `operator<` per comparison.

What if `T` is `std::tuple`?

libc++'s tuple comparison

```
template<size_t _Ip> struct __tuple_less
{
    template <class _Tp, class _Up>
    bool operator() (const _Tp& __x, const _Up& __y)
    {
        return __tuple_less<_Ip-1>() (__x, __y) ||
            (!__tuple_less<_Ip-1>() (__y, __x) && get<_Ip-1>(__x) < get<_Ip-1>(__y));
    }
};

template<> struct __tuple_less<0>
{
    template <class _Tp, class _Up>
    bool operator() (const _Tp&, const _Up&) { return false; }
};

template<class ..._Tp, class ..._Up>
bool operator<(const tuple<_Tp...>& __x, const tuple<_Up...>& __y)
{
    return __tuple_less<sizeof...(_Tp)>() (__x, __y);
}
```

2 tuple comparisons

=

$2n$ element comparisons

So what?

Who uses tuples for anything?

Who *compares* tuples?

Unfortunately, lots of people

<http://vexorian.blogspot.com/2013/07/more-about-c11-tuples-tie-and-maketuple.html>

<http://stackoverflow.com/questions/10806036/using-make-tuple-for-comparison>

<http://stackoverflow.com/questions/6218812/implementing-comparison-operators-via-tuple-and-tie-a-good-idea>

<http://siliconkiwi.blogspot.com/2012/04/stdtie-and-strict-weak-ordering.html>

<http://oraclechang.files.wordpress.com/2013/05/c11-a-cheat-sheete28094alex-sinyakov.pdf>

<http://litedev.wordpress.com/2013/08/12/less-than-obvious/>

<http://wordaligned.org/articles/more-adventures-in-c++>

And on the topic of adding a “spaceship function” to C++:

[Generic compare function \(Adam Badura\)](#)

[Why aren't “tri-valent” comparison functions used in the standard library? \(K. Frank\)](#)

The idiom we want to use in C++14

```
class MyClass {  
    int a, b, c, d;  
public:  
    auto tied() const {  
        return std::tie(a,b,c,d);  
    }  
    bool operator< (const MyClass& rhs) const {  
        return tied() < rhs.tied();  
    }  
};
```

```
... array_pod_sort_comparator<MyClass> ...
```

The idiom we want to use in C++14

```
/// array_pod_sort_comparator - This is helper function for array_pod_sort,  
/// which just uses operator< on T.  
template<typename T>  
int array_pod_sort_comparator (const T& a, const T& b) {  
    if (a < b)  
        return -1;  
    if (b < a)  
        return 1;  
    return 0;  
}
```

```
... array_pod_sort_comparator<MyClass> ...
```

This is disastrously inefficient when `(a == b)`!

Twice as many comparisons as necessary!

How to fix it

We need a trivalent comparison function for tuples.

```
namespace std {  
  
    template <typename T, typename U>  
    int spaceship(const T&, const U&);  
  
    template <typename... T, typename... U>  
    int spaceship(const tuple<T...>&, const tuple<U...>&);  
  
}
```

How to implement it

```
// The easy part.  
//
```

```
namespace std {
```

```
    template <typename T, typename U>  
    int spaceship(const T& x, const U& y)  
    {  
        return (x < y) ? -1 : (y < x) ? 1 : 0;  
    }
```

```
}
```


How to implement it (library style)

```
// The easy part.  
//  
  
namespace std {  
  
    template <class _Tp, class _Up>  
    constexpr int spaceship(const _Tp& __x, const _Up& __y)  
    {  
        return (__x < __y) ? -1 : (__y < __x) ? 1 : 0;  
    }  
  
}
```

How to implement it

```
// The barely harder part.  
//
```

```
namespace std {
```

```
    int spaceship(const string& x, const string& y)  
    {  
        int r = x.compare(y);  
        return (r < 0) ? -1 : (r > 0);  
    }
```

```
}
```

How to implement it (library style)

```
// The barely harder part.  
//  
  
namespace std {  
  
    template<class _Cp, class _Tp, class _Ap,  
              class _Up, class _Bp>  
    int spaceship(const basic_string<_Cp, _Tp, _Ap>& __x,  
                  const basic_string<_Cp, _Up, _Bp>& __y)  
    {  
        int __r = __x.compare(0, __x.size(), __y.data(), __y.size());  
        return (__r < 0) ? -1 : (__r > 0);  
    }  
  
}
```

// The hard part (libc++ style)

```
template<size_t _Ip> struct __tuple_spaceship
{
    template <class _Tp, class _Up>
    constexpr int operator()(const _Tp& __x, const _Up& __y) const
    {
        int __r = __tuple_spaceship<_Ip-1>()(__x, __y);
        return (__r != 0) ? __r : spaceship(get<_Ip-1>(__x), get<_Ip-1>(__y));
    }
};

template<> struct __tuple_spaceship<0>
{
    template <class _Tp, class _Up>
    constexpr int operator()(const _Tp&, const _Up&) const
    {
        return 0;
    }
};

template <class ..._Tp, class ..._Up>
constexpr int spaceship(const tuple<_Tp...>& __x, const tuple<_Up...>& __y)
{
    static_assert(sizeof...(_Tp) == sizeof...(_Up));
    return __tuple_spaceship<sizeof...(_Tp)>()(__x, __y);
}
```

// The hard part (Painless Metaprogramming style)

```
template <class _Tp, class _Up, size_t ..._Ip>
constexpr int __tuple_spaceship(const _Tp& __x, const _Up& __y,
                               const index_sequence<_Ip...>&)
{
    int __r = 0;
    std::initializer_list<int> x = {
        (__r != 0) ? 0 : (__r = spaceship(get<_Ip>(__x), get<_Ip>(__y))) ...
    };
    return __r;
}

template <class ..._Tp, class ..._Up>
constexpr int spaceship(const tuple<_Tp...>& __x, const tuple<_Up...>& __y)
{
    static_assert(sizeof...(_Tp) == sizeof...(_Up), "");
    return __tuple_spaceship(__x, __y, make_index_sequence<sizeof...(_Tp)>{});
}
```

Clang 3.4 doesn't seem to care which way you do it.

Recommended reading: “Towards Painless Metaprogramming” <http://ldionne.github.io/mpl11-cppnow-2014/>

The idiom we *should* use in C++1z

```
class MyClass {
    int a, b, c, d;
public:
    auto tied() const {
        return std::tie(a,b,c,d);
    }
    bool operator< (const MyClass& rhs) const {
        return spaceship(*this, rhs) < 0;
    }
};

int spaceship(const MyClass& a, const MyClass& b) {
    return std::spaceship(a.tied(), b.tied());
}

... array_pod_sort_comparator<MyClass> ...
```

The idiom we *should* use in C++1z

```
/// array_pod_sort_comparator - This is helper function for array_pod_sort,  
/// which just uses "spaceship" on T.  
///  
template<typename T>  
int array_pod_sort_comparator (const T& a, const T& b)  
{  
    using std::spaceship;  
  
    return spaceship (a, b);  
}  
  
... array_pod_sort_comparator<MyClass> ...
```

This is efficient even when `(a == b)`!

Only as many comparisons as necessary!

Unfortunately...

`std::spaceship` is not part of C++1z.

You know anyone on the standards committee?

Postscript

```
struct S
{
    void operator<= (int) {}
};

template<void (S::*)(int)> void f() {}

int main()
{
    f<&S::operator<=>();
}
```

end()