




How Facebook's HHVM Uses C++ for Fun and Profit

September 8th, 2014



Drew Paroski

What is HHVM?

- HHVM = HipHop Virtual Machine
- HHVM is a new open-source virtual machine designed to execute programs written in PHP or Hack*

* Hack is a new language that evolved from PHP which adds a bunch of useful features

What is HHVM?

- Uses trace-based JIT compilation to deliver superior performance while maintaining the flexibility and productivity that PHP developers are used to
- Supports virtually all of the PHP 5.6 language including eval()

HHVM at Facebook

- A lot of core logic for a wide range of Facebook's features and products is written in PHP/Hack
- HHVM was developed to deliver better performance for Facebook's PHP/Hack code base
- HHVM evolved from a PHP->C++ transpiler called "HPHPc"

Performance

- As of mid-2014, HHVM (and HPHPc before it) has realized a $\sim 10x$ increase in throughput and over a 75% reduction in memory usage for facebook.com compared with PHP 5.2

* Notes:

- This compares HHVM circa 2014 with a version of vanilla PHP from 2009
- These figures were calculated by taking multiple delta measurements over multiple years and stitching them together – direct comparison is not possible because FB's codebase is no longer compatible with vanilla PHP
- HHVM/HPHPc and FB's codebase have co-evolved over the past 4 years; until recently HHVM optimization efforts focused solely on FB's codebase

Facebook Data Center



Performance

- The performance of the stock php.net interpreter has notably improved in recent years
- Still, HHVM produces a boost of $\sim 2x$ or more for a lot of PHP applications when compared with PHP 5.5 with opcode cache

Productivity

- HHVM supports the flexible “edit, save, run” development workflow that PHP developers are accustomed to
- When a PHP source file gets modified, HHVM detects that the file changed and recompiles PHP code as needed

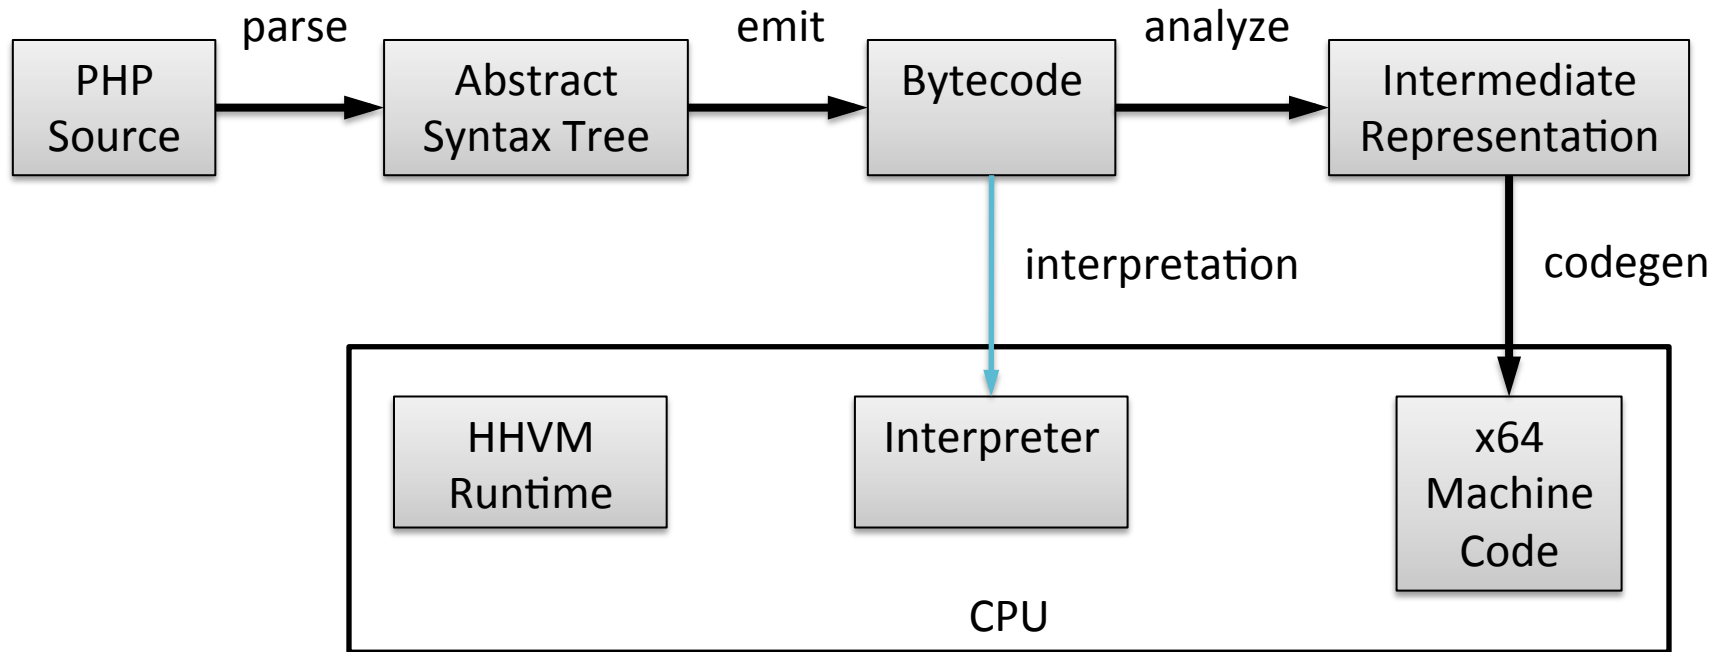
What can HHVM run?

- facebook.com
- Symfony
- MediaWiki
- WordPress
- Laravel
- Magento
- Drupal
- And dozens more of the most popular PHP frameworks and applications

How does HHVM work?

- PHP is parsed and then converted into untyped bytecode and metadata
- Basic blocks of bytecode (tracelets) are analyzed using live type information and converted to a typed IR
- Typed IR is then optimized and converted into x64 machine code

Basic Compilation Pipeline



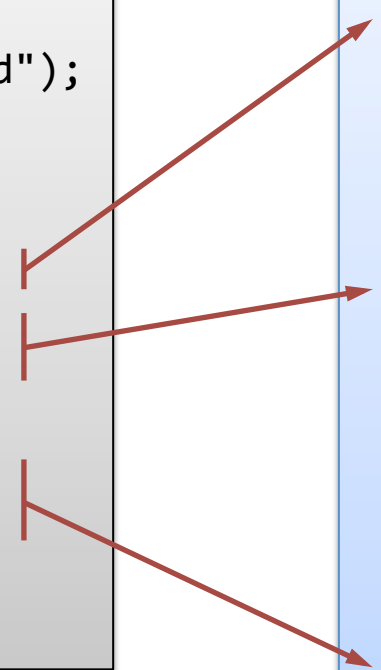
PHP → HHBC compilation

```
1: <?php
2: $a = array("hello" => "world");
3: f($a, 42);
4: f($a, "hello");
5: function f($a, $k) {
6:     if (!empty($a[$k])) {
7:         print "hit: ".$k."\n";
8:         return;
9:     } else {
10:        print "miss: ".$k."\n";
11:        return;
12:    }
13: }
```

```
0: EmptyM <L:0 EL:1>
13: JmpNZ 23 (96)

18: String "hit: "
23: String "\n"
28: CGetL2 1
30: Concat
31: Concat
32: Print
33: PopC
34: Null
35: RetC

36: String "miss: "
41: String "\n"
46: CGetL2 1
48: Concat
49: Concat
50: Print
51: PopC
52: Null
53: RetC
```



Tracelets

- HHVM's JIT compiler translates small amounts of bytecode at a time, using run time type information to drive predictions and generate type-specialized machine code:

Bytecode

```
CGetL 0  
String "x"  
Concat  
CGetL 1  
Eq  
JmpZ 22
```



Live type
info

```
Local 0 is  
type integer  
  
Local 1 is  
type string
```



Type-specialized
translation


Prologue (type guards)
Machine code to perform actual work
Epilogue (jump to the next tracelet)

Translation example

```
function f($a, $k) {  
    if (!empty($a[$k])) {  
        echo "hit $k\n";  
        return;  
    } else {  
        echo "miss $k\n";  
        return;  
    }  
}
```

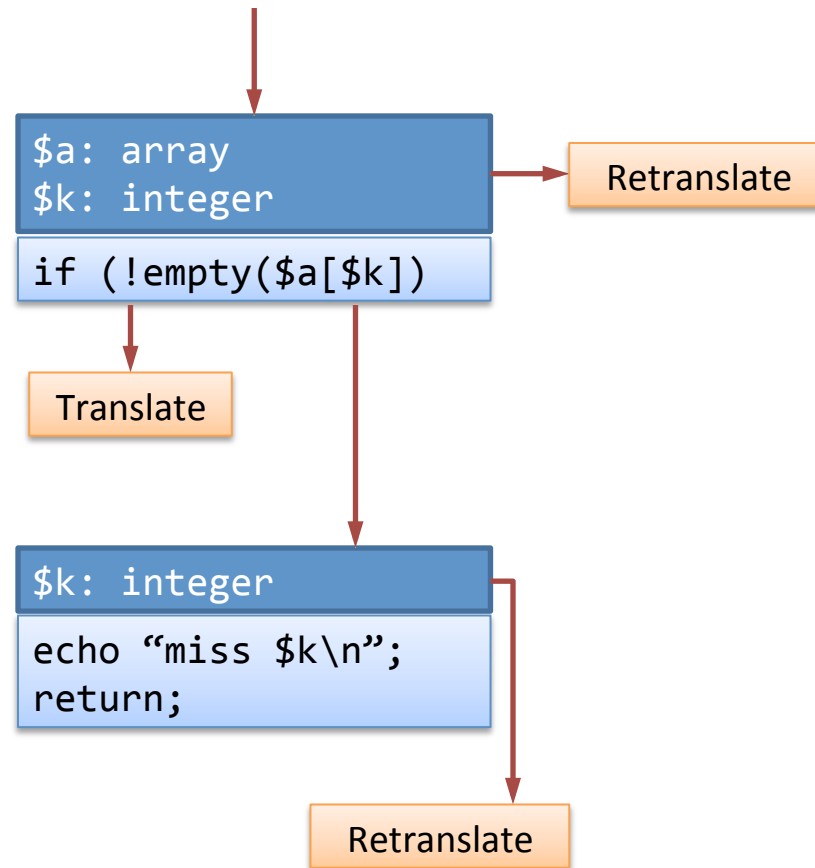
Translation example

```
function f($a, $k) {  
    if (!empty($a[$k])) {  
        echo "hit $k\n";  
        return;  
    } else {  
        echo "miss $k\n";  
        return;  
    }  
}
```



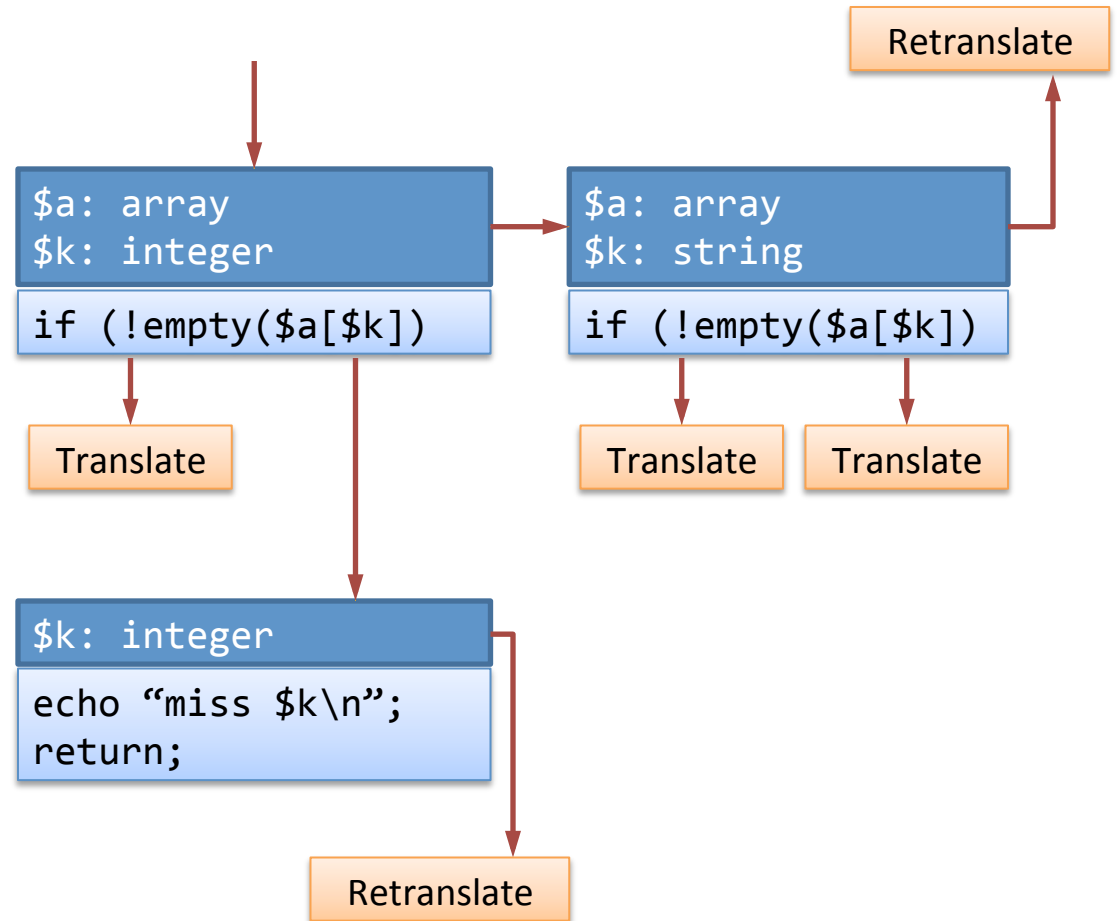
Translation example

```
function f($a, $k) {  
    if (!empty($a[$k])) {  
        echo "hit $k\n";  
        return;  
    } else {  
        echo "miss $k\n";  
        return;  
    }  
}
```



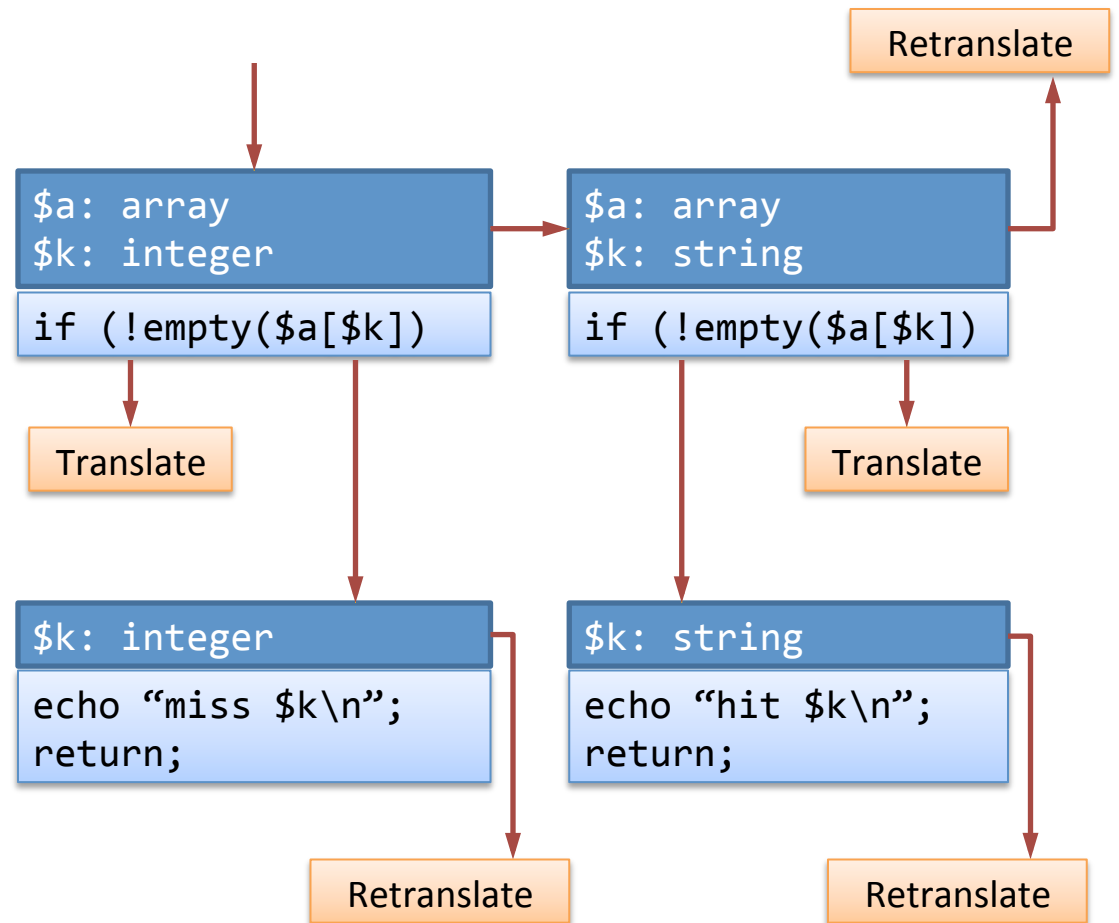
Translation example

```
function f($a, $k) {  
    if (!empty($a[$k])) {  
        echo "hit $k\n";  
        return;  
    } else {  
        echo "miss $k\n";  
        return;  
    }  
}
```



Translation example

```
function f($a, $k) {  
    if (!empty($a[$k])) {  
        echo "hit $k\n";  
        return;  
    } else {  
        echo "miss $k\n";  
        return;  
    }  
}
```



Another translation example

```
$n = 3 * $n + 1;  
...
```



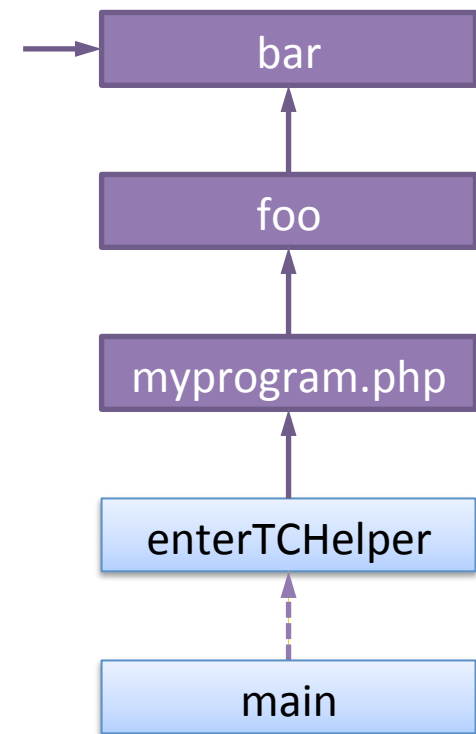
```
Int 3  
CGetL 0  
Mul  
Int 1  
Add  
SetL 0  
PopC  
...
```



```
cmpl    $0xa, -0x8(%rbp)  
jne     __retranslate  
mov     -0x10(%rbp), %rax  
mov     %rax, %rcx  
shl     %rcx  
add     %rcx, %rax  
mov     $0x1, %r13d  
add     %rax, %r13  
...
```

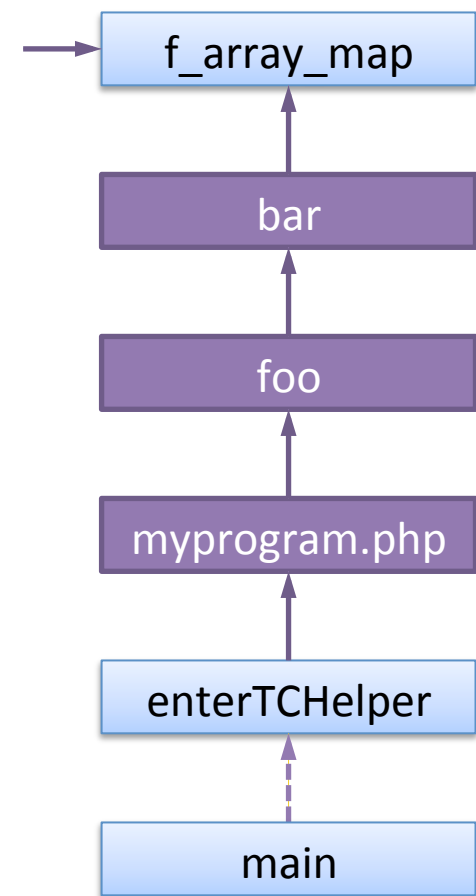
HHVM's Execution Model

- HHVM models the flow of execution using a stack of frames referred to as the ***call stack***
- Each frame represents a function invocation
- VM frames correspond to PHP function invocations (shown in purple)
- C++ frames correspond to C++ function invocations (shown in blue)



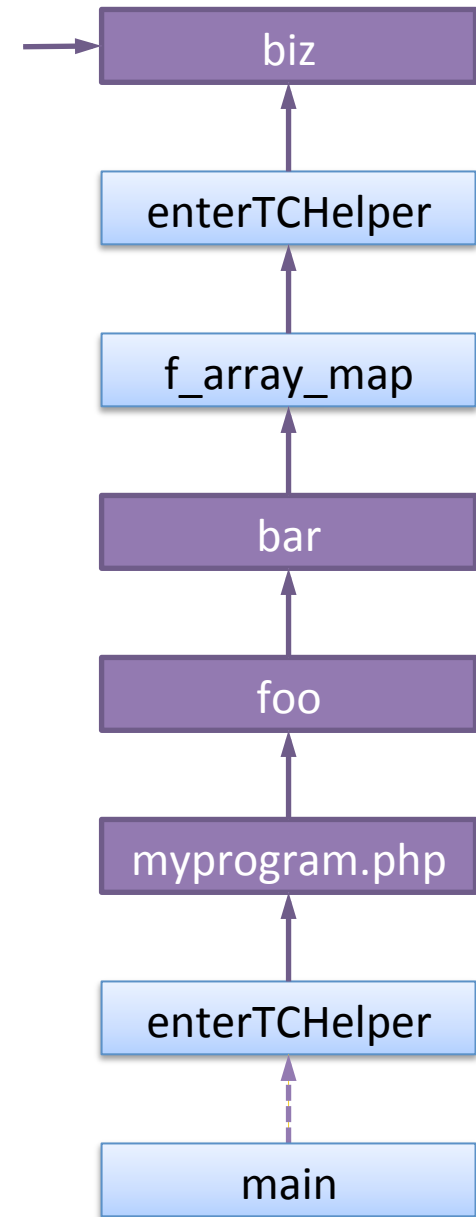
HHVM's Execution Model

- HHVM models the flow of execution using a stack of frames referred to as the ***call stack***
- Each frame represents a function invocation
- VM frames correspond to PHP function invocations (shown in purple)
- C++ frames correspond to C++ function invocations (shown in blue)
- PHP code can call into C++ code



HHVM's Execution Model


- HHVM models the flow of execution using a stack of frames referred to as the ***call stack***
- Each frame represents a function invocation
- VM frames correspond to PHP function invocations (shown in purple)
- C++ frames correspond to C++ function invocations (shown in blue)
- PHP code can call into C++ code
- C++ code can call back into PHP code



Why was C++ a good choice
for building HHVM?

Why was C++ a good choice for building HHVM?

- C++ hits a sweet spot
 - Superb control and performance on one end
 - Convenience, maintainability, and safety on the other
 - Gives the programmer fine-grained control to choose between these competing concerns as desired in different parts of the codebase



Why was C++ a good choice for building HHVM?

- Performance is equal or better than any other language (aside from assembly)
- Gives the programmer an incredible amount of freedom and control:
 - Manual memory management
 - Unsafe casting, field size, and layout
 - Flexible, light-weight interop with assembly and machine code

Why was C++ a good choice for building HHVM?

- C++ offers many convenient features that can be used as little or as much as desired
 - Virtual methods
 - Multiple inheritance
 - reinterpret_cast vs. dynamic_cast
 - Plain old data vs. constructors/destructors
 - Raw pointers vs. references vs. smart pointers
 - Stack allocation vs. malloc vs. new
 - Templates and macros

Templates

- Templates are great for maintainability
- More hygienic than preprocessor macros
- For HHVM, templates were particularly useful for critical parts of the engine where we wanted the compiler to do as much inlining as possible to improve perf

Templates

- Eliminating branches:

```
void foo(bool b, ..) {  
    if (b) {  
        bar();  
    }  
    ..  
}
```



```
test %rdi, %rdi  
jz L1  
call 0x4005a0 <bar>  
L1:
```

```
template <bool b>  
void foo(..) {  
    if (b) {  
        bar();  
    }  
    ..  
}
```



```
call 0x4005a0 <bar>
```

Templates

- Reducing indirection:

```
void foo(int(*fn)(), ..)  
{  
    int x = fn();  
    ..  
}
```



```
call *%rdi
```

```
template <class T>  
void foo(..) {  
    int x = T::staticMeth();  
    ..  
}
```



```
call 0x400780 <Foo::staticMeth>
```

C++11's Lambdas

- Lambdas are useful because they help keep related pieces of logic together in one place
- For HHVM, we typically use reference capture (i.e. “[&]”) and we’re careful about making sure lambdas do not outlive the captured variables on the stack

C++11's Lambdas

```
void CodeGenerator::cgCountArray(IRInstruction* inst) {  
    ...  
    ifThenElse(vmain(), vcold(), CC_Z,  
        [&](Vout& v) {  
            cgCallNative(v, inst);  
        },  
        [&](Vout& v) {  
            v << loadl{baseReg[ArrayData::offsetofSize()],  
                dstReg};  
        }  
    );  
    ...  
}
```

X Macros

- C++'s preprocessor is unhygienic, but it is extremely powerful and it can be very useful
- The ***X Macro*** technique can help with maintainability if used judiciously
- HHVM's bytecode definitions use the X Macro technique to make it easy to add, remove, or modify bytecode instructions

X Macros

```
#define OPCODES \  
    O(PopC,    NA,          ONE(CV),      NOV,      NF) \  
    O(PopV,    NA,          ONE(VV),      NOV      NF) \  
    O(CGetL,   ONE(LA),     NOV,          ONE(CV),   NF) \  
    O(Add,     NA,          TWO(CV,CV),    ONE(CV),   NF) \  
    ..  
  
enum class Op : uint8_t {  
#define O(name, ...) name,  
    OPCODES  
#undef O  
};
```

X Macros

```
#define OPCODES \  
    O(PopC,    NA,          ONE(CV),      NOV,          NF) \  
    O(PopV,    NA,          ONE(VV),      NOV           NF) \  
    O(CGetL,   ONE(LA),     NOV,           ONE(CV),      NF) \  
    O(Add,     NA,          TWO(CV,CV),    ONE(CV),      NF) \  
    ..  
#define PUSH_NOV                /* nop */  
#define PUSH_ONE(t)              PUSH_##t  
#define PUSH_TWO(t1, t2)        PUSH_##t2; PUSH_##t1  
#define PUSH_CV                  ..  
#define PUSH_VV                  ..  
..  
#define O(name, imm, push, pop, flags) \  
    .. PUSH_##push ..  
OPCODES  
#undef O  
#undef PUSH_NOV  
..
```

Unions and field size

- Unions are super useful when dealing with dynamically-typed values
- Unions can also be used to reduce the size of structs that have mutually exclusive fields
- Ability to control field size also comes in handy:
 - 8-bit, 16-bit, 32-bit, or 64-bit integers

TypedValue union

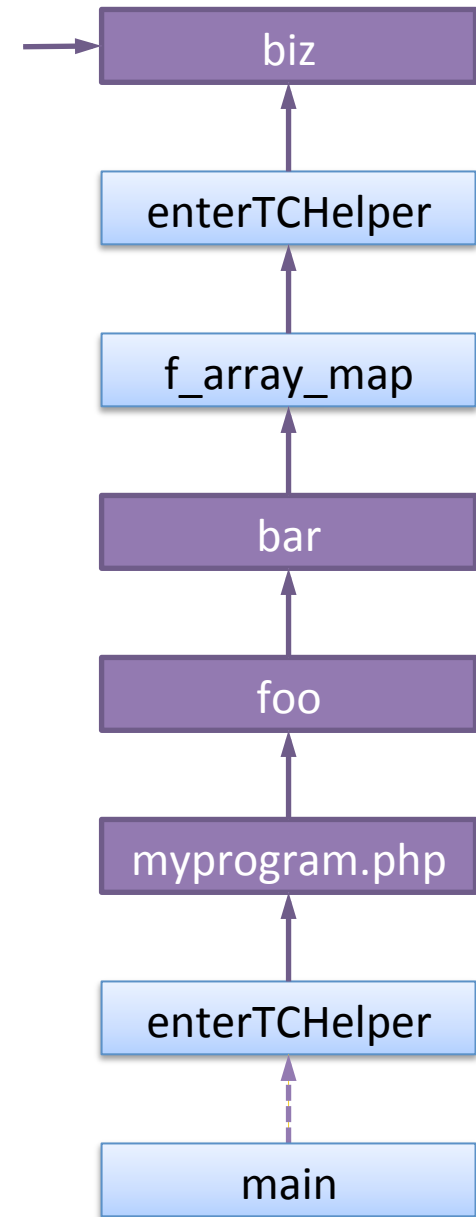
```
struct TypedValue {  
    union {  
        int64_t      num;  
        double       dbl;  
        StringData*  pstr;  
        ArrayData*   parr;  
        ObjectData*  pObj;  
        ResourceData* pres;  
        Class*       pcls;  
        RefData*     pref;  
    } m_data;  
    DataType m_type;  
    AuxUnion m_aux;  
};  
  
enum DataType : int8_t {  
    KindOfClass      = -13,  
    KindOfUninit     = 0x00,  
    KindOfNull       = 0x08,  
    KindOfBoolean    = 0x09,  
    KindOfInt64      = 0x0a,  
    KindOfDouble     = 0x0b,  
    KindOfString     = 0x14,  
    KindOfArray      = 0x20,  
    KindOfObject     = 0x30,  
    KindOfResource   = 0x40,  
    KindOfRef        = 0x50,  
    ..  
};
```

Unsafe casts and bit-stealing

- C++ allows for unsafe casts between integers and different pointer types
- The implementation of malloc used by HHVM always returns chunks of memory aligned to 8-byte boundaries
- Unsafe casts and bit masking can be used to steal the low bits of pointers

Activation Records

- HHVM's `ActRec` struct is used as the header for each VM frame
- `ActRecs` store essential information about the PHP function invocation, such as:
 - The name and other metadata pertaining to the current function
 - Where to jump to when the function returns
 - Necessary bookkeeping to support getting a PHP backtrace (i.e. `debug_backtrace()`)



Bit-stealing with ActRecs

```
struct ActRec {  
    ActRec* m_savedFp;  
    uint64_t m_savedRip;  
    Func* m_func;  
    uint32_t m_soff;  
    uint32_t m_numArgsAndFlags;  
    union {  
        ObjectData* m_this;  
        Class* m_cls;  
    };  
    union {  
        VarEnv* m_varEnv;  
        ExtraArgs* m_extraArgs;  
        StringData* m_invName;  
    };  
};
```

Saved FP		Saved RIP	
Func		Bytecode offset	# args / flags
		VarEnv / ExtraArgs / InvName	

Bit-stealing with ActRecs

```
struct ActRec {  
    ActRec* m_savedFp;  
    uint64_t m_savedRip;  
    Func* m_func;  
    uint32_t m_soff;  
    uint32_t m_numArgsAndFlags;  
    union {  
        ObjectData* m_this;  
        Class* m_cls;  
    };  
    union {  
        VarEnv* m_varEnv;  
        ExtraArgs* m_extraArgs;  
        StringData* m_invName;  
    };  
};
```

Saved FP		Saved RIP	
Func		Bytecode offset	# args / flags
"this" pointer / late-bound class		VarEnv / ExtraArgs / InvName	

Bit-stealing with ActRecs

- For unions of pointers, we steal the low bit(s) and use it as a tag to disambiguate

```
bool hasThis() const {  
    return m_this && !((intptr_t)m_this & 1);  
}  
ObjectData* getThis() {  
    return m_this;  
}  
void setThis(ObjectData* val) {  
    m_this = val;  
}  
bool hasClass() {  
    return (intptr_t)m_cls & 1;  
}  
Class* getClass() {  
    return (Class*)((intptr_t)m_cls & 1);  
}  
void setClass(Class* val) {  
    m_cls = (Class*)((intptr_t)val | 1);  
}
```

Bit-stealing with ActRecs

```
struct ActRec {  
    ActRec* m_savedFp;  
    uint64_t m_savedRip;  
    Func* m_func;  
    uint32_t m_soff;  
    uint32_t m_numArgsAndFlags;  
    union {  
        ObjectData* m_this;  
        Class* m_cls;  
    };  
    union {  
        VarEnv* m_varEnv;  
        ExtraArgs* m_extraArgs;  
        StringData* m_invName;  
    };  
};
```

Saved FP		Saved RIP	
Func		Bytecode offset	# args / flags
"this" pointer / late-bound class		VarEnv / ExtraArgs / InvName	

Bit-stealing with ActRecs

- For `m_numArgsAndFlags`, we steal the high bit(s) for flags

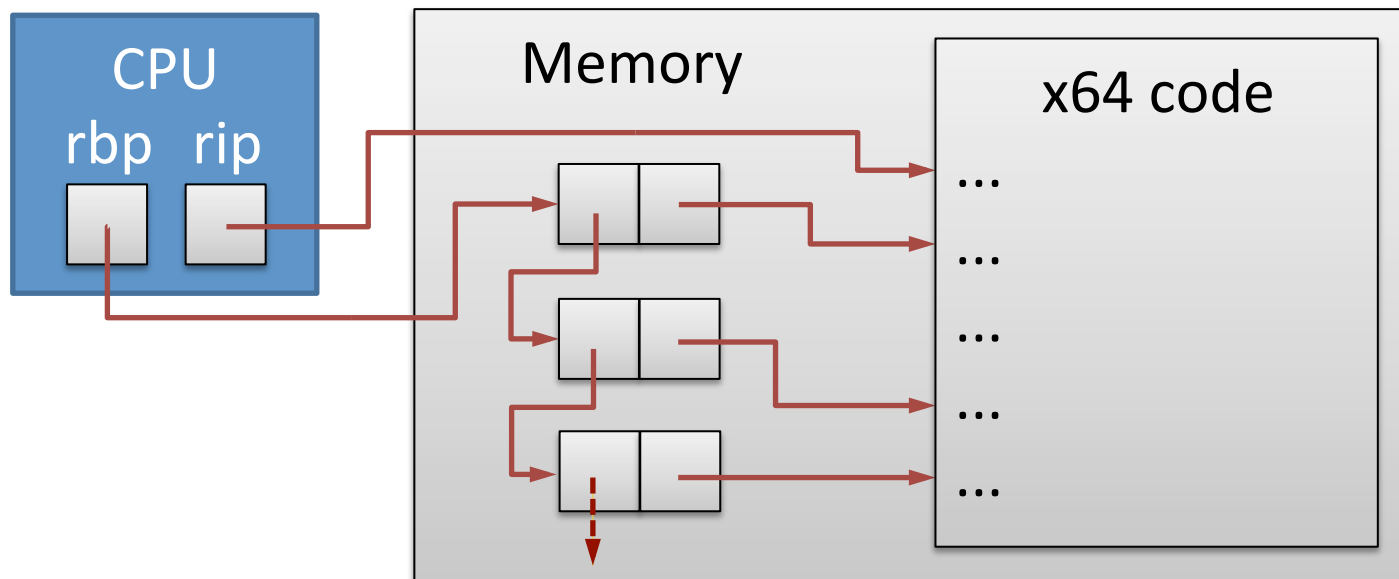
```
int32_t numArgs() {  
    return m_numArgsAndFlags & 0x7fffffff;  
}  
bool isCtorFrame() {  
    return m_numArgsAndFlags & (1 << 31);  
}  
void initNumArgs(uint32_t numArgs) {  
    assert(!(numArgs & (1 << 31)));  
    m_numArgsAndFlags = numArgs;  
}  
void initNumArgsCtorFrame(uint32_t numArgs) {  
    assert(!(numArgs & (1 << 31)));  
    m_numArgsAndFlags = numArgs | (1 << 31);  
}
```

Profiling

- The C++ ecosystem has lots of mature tools for profiling
- Linux's `perf` tool is awesome
 - Uses a sampling technique to determine which functions a program spends the most CPU time on
 - Keeps track of the full call stack for each sample

Profiling

- On x64, Linux's `perf` tool works by using the ***rbp-chain*** to walk the call stack
- C++ code needs be compiled with frame pointers (i.e. `-fno-omit-frame-pointer`)



Profiling

- ActRec is designed so the first 16 bytes has the same layout as C++ frames
- HHVM sets up VM frames so that they participate in the rbp-chain
- This makes it possible to profile PHP programs running under HHVM using Linux's `perf` tool

Saved FP		Saved RIP	
Func		Bytecode offset	# args / flags
This / Cls		VarEnv / ExtraArgs / InvName	

Events: 4K cycles

```

- 1.02% hhvm hhvm      [.] HPHP::ObjectData::setProp(HPHP::Class*, HPHP::S
- HPHP::ObjectData::setProp(HPHP::Class*, HPHP::StringData const*, HPHP::Typed
  - 73.64% HPHP::JIT::MInstrHelpers::setPropC0(HPHP::Class*, HPHP::TypedValue
    + 57.02% PHP::/www/wordpress/wp-includes/post.php::WP_Post::__construct
    + 16.74% PHP::/www/wordpress/wp-includes/post.php::WP_Post::__construct
    + 16.61% PHP::/www/wordpress/wp-includes/post.php::sanitize_post
    + 16.47% HPHP::ObjectData::o_setArray(HPHP::Array const&)
    + 3.98% PHP::/www/wordpress/wp-includes/post.php::WP_Post::__construct
- 0.80% hhvm hhvm      [.] HPHP::f_in_array(HPHP::Variant const&, HPHP::Va
- HPHP::f_in_array(HPHP::Variant const&, HPHP::Variant const&, bool)
  + 49.94% PHP::/www/wordpress/wp-includes/post.php::sanitize_post_field
  + 26.36% PHP::/www/wordpress/wp-includes/post.php::sanitize_post_field
  + 5.86% PHP::/www/wordpress/wp-includes/option.php::get_option
  + 3.90% PHP::/www/wordpress/wp-includes/post.php::sanitize_post
- 0.73% hhvm hhvm      [.] HPHP::Unit::GetNamedEntity(HPHP::StringData con
- HPHP::Unit::GetNamedEntity(HPHP::StringData const*, bool, HPHP::String*)
  - 31.10% HPHP::f_is_a(HPHP::Variant const&, HPHP::String const&, bool)
    + 73.09% PHP::/www/wordpress/wp-includes/post.php::get_post
    + 25.37% PHP::/www/wordpress/wp-includes/post.php::get_post
    + 1.17% PHP::/www/wordpress/wp-includes/post.php::get_post
  - 27.98% HPHP::Unit::loadFunc(HPHP::StringData const*)
    - 85.99% HPHP::JIT::fpushCufHelperString(HPHP::StringData*, HPHP::ActRec
      + 97.73% PHP::/www/wordpress/wp-includes/plugin.php::apply_filters
      + 0.77% PHP::/www/wordpress/wp-includes/plugin.php::apply_filters

```

Press '?' for help on key bindings

Exception handling / unwinding

- `setjmp` / `longjmp` can be useful when implementing exception handling for a VM
- HHVM's initial EH implementation used `setjmp` and `longjmp` to skip over the `enterTCHelper` trampoline and VM frames when an exception was thrown
 - This scheme was a bit clunky but it got HHVM's EH system up and running quickly

Exception handling / unwinding

- Later, HHVM switched over to using g++'s `__register_frame()` function to integrate with the C++ runtime's exception unwinding system by registering a “personality routine”

```
_Unwind_Reason_Code  
tc_unwind_personality(  
    int version,  
    _Unwind_Action actions,  
    uint64_t exceptionClass,  
    _Unwind_Exception* exceptionObj,  
    _Unwind_Context* context  
);
```

C++ / machine code interop

- On x64 and other popular platforms, C++ has a well-defined ABI for function calls
- Generating calls from machine code to C++ functions is light-weight and easy^{*}

* Provided that parameters are pointers or primitive types and virtual methods are not involved

C++ / machine code interop

- Making C++ call into machine code is also relatively simple
- C++ allows the programmer to take a `void*` that points to machine code and cast it to a function pointer and invoke it:

```
void* p = getMachineCodeAddress();  
typedef int(*FuncPtr)(int);  
FuncPtr fn = (FuncPtr)p;  
int result = fn(123);
```

enterTCHelper example

enterTCHelper:

```
push %rbp
push %rcx
mov %rdi,%rbx
mov %r8,%r12
mov %rsi,%rbp
call *%rdx
pop %rbx
mov %rdi,0x0(%rbx)
mov %rsi,0x8(%rbx)
...
pop %rbp
ret
```

// Parameters: rdi, rsi, rdx, rcx, r8

extern "C" void

```
enterTCHelper(Cell* sp, ActRec* fp,
               TCA start, TReqInfo* info,
               void* tlBase);
```

inline void


```
enterTC(TCA start, TReqInfo& info) {
    asm volatile("" : : :
                 "rbx", "r12", "r13", "r14", "r15");
    auto& regs = vmRegsUnsafe();
    enterTCHelper(regs.stkTop(), regs.fp,
                  start, &info, tl_base());
    asm volatile("" : : :
                 "rbx", "r12", "r13", "r14", "r15");
}
```

Manual Memory Management

- C++ gives the programmer freedom to manually manage memory; this was essential for HHVM
- High compatibility with the php.net interpreter was important; given PHP's semantics, refcounting was less risky
 - Thus tracing GC was not an attractive option at the outset

Native allocation

- C++ provides an easy and light-weight means for the programmer to choose which implementation of malloc they want to use




jemalloc

- HHVM uses jemalloc to handle calls to malloc APIs such as `malloc()`, `free()`, `calloc()`, etc.

jemalloc

- jemalloc is a world-class concurrent memory allocator used by the FreeBSD operating system, Firefox, and many other software projects



jemalloc

- Very efficient both for single- and multi-threaded programs
 - Particularly good at minimizing fragmentation for long-running processes

jemalloc

- HHVM takes advantage of jemalloc-specific APIs that aren't part of the standard malloc interface
- Some features were added to jemalloc to specifically to help out HHVM
 - jemalloc's `mallocx()` and `allocm()` APIs were updated to support allocating memory with “low” addresses that can fit within 32 bits
 - Sophisticated heap profiling functionality was added to jemalloc to aid with investigating how to improve HHVM's performance

Low addresses

- Allocating memory with low addresses helps make HHVM's data structures smaller and helps with generating more efficient machine code

```
struct LowClassPtr {  
    int32_t m_raw;  
    Class* get() {  
        return reinterpret_cast<Class*>(m_raw);  
    }  
    void set(Class* c) {  
        m_raw = reinterpret_cast<int32_t>(c);  
    }  
};
```

Low addresses

```
movl $0x7671a20, -0x20(%rbx)
```

```
[c7 43 e0 20 1a 67 07]
```

7 bytes total

```
movabs $0x7fffe8408000, %rax
```

```
mov     %rax, -0x20(%rbx)
```

```
[48 b8 00 80 40 e8 ff 7f 00 00 48 89 43 e0]
```

14 bytes total

Memory Management for PHP

- The PHP language requires that engines provide some form of automatic memory management for programs written in PHP
- PHP models concurrency using separate *requests*
- Each request has its own distinct heap, and at the end of a request the entire heap dies
- A given request cannot directly access the heap of another request

HHVM's request allocator

- HHVM implements its own custom ***request allocator***
 - Maintains separate isolated heap for each request; this means refcounting doesn't need to use atomic inc/dec
 - Optimizes reclamation at the end of the request
 - Avoids some of the overheads of malloc
 - Controlling the allocator implementation makes it easier to optimize how JIT'd code interacts with the allocator

Huge pages

- Manual memory management also made it possible for HHVM to take advantage of `madvise()`'s `MADV_HUGEPAGE` feature which takes advantage of larger page sizes supported by the TLB hardware
 - By using fewer iTLB entries we were able to significantly reduce iTLB misses, which gave us a nice boost for larger PHP codebases

Obstacles

- There were some features of C++ that posed challenges for us when building HHVM
- For most of these obstacles there was a way to hack around them, which is a testament to C++'s power and flexibility

Unions and non-POD types

- Before C++11, unions could only work with “plain old data” (POD) types
- For HHVM, we wanted to reuse parts of HPHPC’s runtime that dealt with iterating over PHP arrays, but we encountered a problem where we needed to make a union of non-POD types

Unions and non-POD types

```
struct Iter {  
    bool init(TypedValue* c) {  
        new (&arr()) ArrayIter(c);  
    }  
    bool minit(TypedValue* v) {  
        new (&marr()) MArrayIter(v);  
    }  
    void free() { arr().~ArrayIter(); }  
    void mfree() { marr().~MArrayIter(); }  
    ArrayIter& arr() { return *(ArrayIter*)m_u; }  
    MArrayIter& marr() { return *(MArrayIter*)m_u; }  
    char m_u[MAX(sizeof(ArrayIter),  
                  sizeof(MArrayIter))];  
} __attribute__((aligned(16)));
```


C++11's Unrestricted Unions

```
struct Iter {  
    bool init(TypedValue* c) {  
        new (&m_u.arr) ArrayIter(c);  
    }  
    bool minit(TypedValue* v) {  
        new (&m_u.marr) ArrayIter(v);  
    }  
    void free() { m_u.arr.~ArrayIter(); }  
    void mfree() { m_u.marr.~MArrayIter(); }  
    union Data {  
        ArrayIter arr;  
        MArrayIter marr;  
    } m_u;  
}
```

Unnecessary refcounting with smart pointers

- Before move constructors and rvalues were introduced in C++11, smart pointers would often do unnecessary refcounting:

```
class Variant {  
    Variant(String s) { .. }  
};  
class String {  
    ..  
};  
String foo() { .. }
```

```
Variant v(foo());  
..
```

Avoiding unnecessary refcounting

- Prior to C++11, we did some awkward dances to avoid unnecessary refcounting:

```
class Variant {  
    enum NoInc { noInc = 0 };  
    Variant(StringData* s, NoInc) { .. }  
};  
class String {  
    StringData* detach() {  
        auto p = m_data;  
        m_data = nullptr;  
        return p;  
    }  
};  
String foo() { .. }  
  
Variant v(foo().detach(), Variant::noInc);  
..
```

C++11's Move Constructors

- Move constructors made it a lot easier to avoid unnecessary refcounting with smart pointers:

```
class Variant {  
    Variant(String&& s) { .. s.detach() .. }  
};  
class String {  
    StringData* detach() {  
        auto p = m_data;  
        m_data = nullptr;  
        return p;  
    }  
};  
String foo() { .. }  
  
Variant v(foo());  
..
```

C++ instance methods

- Generating machine code that calls into a non-virtual C++ instance method was a bit difficult
 - Getting at the machine code address for the method not straight-forward
- For HHVM, we wanted to reuse some existing parts of HPPc's runtime that used non-virtual C++ instance methods
- We found a way to make it work for g++

C++ instance methods

```
template <typename MethPtr>
void* getMethodPtr(MethPtr p) {
    union U { MethPtr meth; void* ptr; };
    return ((U*)&p)->ptr;
}
class C {
    public: int foo(int x) { .. }
};
void generateCallSite() {
    void* addr = getMethodPtr(&C::foo);
    printf("callq 0x%x\n", addr);
}
```

// Example output

```
callq 0x400570
```

C++ instance methods

```
template <typename MethPtr>
void* getMethodPtr(MethPtr p) {
    union U { MethPtr meth; void* ptr; };
    return ((U*)&p)->ptr;
}
class C {
    public: int foo(int x) { .. }
};
void test(C* c) {
    typedef int (*FuncPtr)(void*,int);
    void* addr = getMethodPtr(&C::foo);
    ((FuncPtr)addr)(c, 123);
}
```

C++ virtual methods

- Generating machine code that calls into a C++ virtual method was tricky
- We found a way to make it work for g++
- For HHVM, we wanted to reuse some existing parts of HPPc's runtime that used C++ virtual methods

C++ virtual methods

```
template <typename MethPtr>
int getVTableOffset(MethPtr p) {
    union U { MethPtr meth; int64_t off; };
    return ((U*)&p)->off - 1;
}

class C {
    public: virtual int foo(int x) { .. }
};

void generateCallSite() {
    int off = getVTableOffset(&C::foo);
    printf("mov (%rdi), %rax\n");
    printf("callq *0x%x(%rax)\n", (int)off);
}
```

// Example output

```
mov (%rdi), %rax
callq *0x8(%rax)
```

C++ virtual methods

```
template <typename MethPtr>
int getVTableOffset(MethPtr p) {
    union U { MethPtr meth; int64_t off; };
    return ((U*)&p)->off - 1;
}

void* getVirtMethAddr(void* obj, int off) {
    return *(void**)(*(intptr_t*)obj + off);
}

class C {
    public: virtual int foo(int x) { .. }
};

void test(C* c) {
    typedef int (*FuncPtr)(void*,int);
    int64_t off = getVTableOffset(&C::foo);
    ((FuncPtr)getVirtMethAddr(c, off))(c, 123);
}
```



Questions?

Website:	hhvm.com
Facebook Page:	facebook.com/hhvm
Github:	github.com/facebook/hhvm