

# Parallelizing the Standard Algorithms Library

Jared Hoberock

NVIDIA

Programming Systems and Applications Research Group

CppCon

September, 2014

# Bringing Parallelism to C++

Technical Specification for Parallel Algorithms

Multi-vendor collaboration

Based on proven technologies

- Thrust (NVIDIA)
- PPL (Microsoft)
- TBB (Intel)

Multiple implementations in progress

Targeting C++17

# Roadmap

Parallelism?

Motivating example

What's included in the box

The details

Future work

# What do I mean by parallelism?

That's like threads, right?

When I say “parallel”, think “independent”


- Concurrency is an optimization
- Concurrency can be hard
- locking, exclusion, communication, shared state, data races...

# What do I mean by parallelism?

Parallel programming => identifying tasks which may be performed independently

How to communicate this information to the system?

It's easy!



# Simple parallelism for everyone

Easy to access

Interoperability with existing codes

Supported as broadly as possible



Concurrency is an invisible optimization

Vendor extensible

# Motivating Example

# Motivating Example: Weld Vertices

Marching Cubes Algorithm






# Motivating Example: Weld Vertices

Problem: Marching Cubes produces “triangle soup”

Solution: “Weld” redundant vertices together



# Motivating Example: Weld Vertices

Easy with the right high-level algorithms

Procedure:

1. Sort triangle vertices
2. Collapse spans of like vertices
3. Search for each vertex's unique index

# Motivating Example: Weld Vertices

```
using namespace std;

using vertex = tuple<float,float>;
vector<vertex> vertices = input;
vector<size_t> indices(input.size());

// sort vertices to bring duplicates together
sort(vertices.begin(), vertices.end());

// find unique vertices and erase redundancies
auto redundant_begin = unique(vertices.begin(), vertices.end());
vertices.erase(redundant_begin, vertices.end());

// find index of each vertex in the list of unique vertices
my_find_all_lower_bounds(vertices.begin(), vertices.end(),
                        input.begin(), input.end(),
                        indices.begin());
```

**Now do it in parallel?**

# Easy!

```
using namespace std;
using namespace std::experimental::parallel;

using vertex = tuple<float,float>;
vector<vertex> vertices = input;
vector<size_t> indices(input.size());

// sort vertices to bring duplicates together
sort(par, vertices.begin(), vertices.end());

// find unique vertices and erase redundancies
auto redundant_begin = unique(par, vertices.begin(), vertices.end());
vertices.erase(redundant_begin, vertices.end());

// find index of each vertex in the list of unique vertices
my_find_all_lower_bounds(par, vertices.begin(), vertices.end(),
                        input.begin(), input.end(),
                        indices.begin());
```

# Wait, I changed my mind...

```
using namespace std;
using namespace std::experimental::parallel;

using vertex = tuple<float,float>;
vector<vertex> vertices = input;
vector<size_t> indices(input.size());

// sort vertices to bring duplicates together
sort(seq, vertices.begin(), vertices.end());

// find unique vertices and erase redundancies
auto redundant_begin = unique(seq, vertices.begin(), vertices.end());
vertices.erase(redundant_begin, vertices.end());

// find index of each vertex in the list of unique vertices
my_find_all_lower_bounds(seq, vertices.begin(), vertices.end(),
                        input.begin(), input.end(),
                        indices.begin());
```

# Don't make me choose!

```
using namespace std;
using namespace std::experimental::parallel;

...

execution_policy exec = seq;
if(input.size() > some_threshold) exec = par;

// sort vertices to bring duplicates together
sort(exec, vertices.begin(), vertices.end());

// find unique vertices and erase redundancies
auto redundant_begin = unique(exec, vertices.begin(), vertices.end());
vertices.erase(redundant_begin, vertices.end());

// find index of each vertex in the list of unique vertices
my_find_all_lower_bounds(exec, vertices.begin(), vertices.end(),
                        input.begin(), input.end(),
                        indices.begin());
```

# my\_find\_all\_lower\_bounds

```
template<class ExecutionPolicy,
         class ForwardIterator,
         class InputIterator,
         class OutputIterator>
OutputIterator my_find_all_lower_bounds(ExecutionPolicy& exec,
                                       ForwardIterator haystack_begin,
                                       ForwardIterator haystack_end,
                                       InputIterator needles_begin,
                                       InputIterator needles_end,
                                       OutputIterator result)
{
    return transform(exec, needles_begin, needles_end, result,
                    [=](auto& needle)
                    {
                        auto iter = std::lower_bound(haystack_begin, haystack_end, needle);
                        return std::distance(haystack_begin, iter);
                    });
}
```



# my\_find\_all\_lower\_bounds

Truly general

- generic in data types (via iterators)
- generic in execution (via execution policy)

Composing our higher-level algorithms from lower-level primitives gives us parallelism for free!

# How to write parallel programs

High-level algorithms

Control sequential/parallel execution with policies

Communicate dependencies

# What's included in the box

Execution Policies

Parallel Algorithms

Parallel Exceptions

# Execution Policies

```
using namespace std::experimental::parallelism;
std::vector<int> data = ...

// vanilla sort
sort(data.begin(), data.end());

// explicitly sequential sort
sort(seq, data.begin(), data.end());

// permitting parallel execution
sort(par, data.begin(), data.end());

// permitting vectorization as well
sort(par_vec, data.begin(), data.end());
```

# Execution Policies

```
// sort with dynamically-selected execution
size_t threshold = ...
execution_policy = seq;
if(vec.size() > threshold)
{
    exec = par;
}

sort(exec, vec.begin(), vec.end());
```

# Execution Policies

```
sort(vectorize_in_this_thread, vec.begin(), vec.end());
```

```
sort(submit_to_my_thread_pool, vec.begin(), vec.end());
```

```
sort(execute_on_that_gpu, vec.begin(), vec.end());
```

```
sort(offload_to_my_fpga, vec.begin(), vec.end());
```

```
sort(send_to_the_cloud, vec.begin(), vec.end());
```

```
sort(laundry_through_botnet, vec.begin(), vec.end());
```

Implementation-Defined  
(Non-Standard)

# What is an execution policy?

Promise that a particular kind of reordering will preserve the meaning of a program

Request that the implementation use some sort of execution strategy

What sorts of reordering are allowable?


# The Details



# Sequential Execution

```
algo(seq, begin, end, func);
```

`algo` invokes function  
and iterator operations  
in sequential order in  
the calling thread



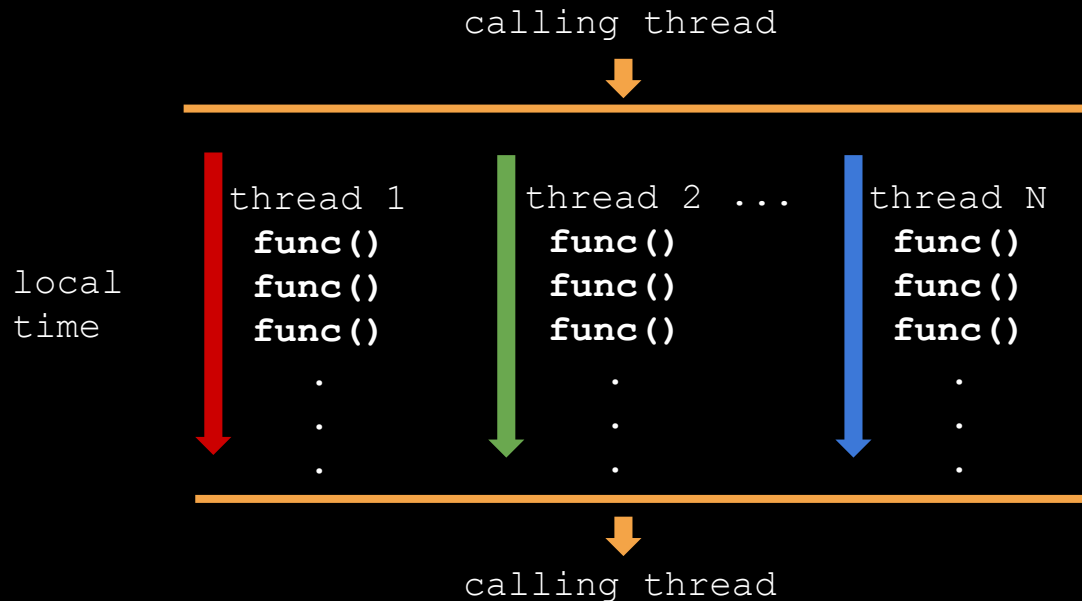
# Parallelizable Execution

```
algo(par, begin, end, func);
```

`algo` is permitted to invoke function  
unsequenced if invoked in different threads,  
and indeterminate order if invoked in the same  
thread

# Parallelizable Execution

```
algo(par, begin, end, func);
```



# Parallelizable Execution

It is the caller's responsibility to ensure correctness, for example that the invocation does not introduce data races or deadlocks.

# Parallelizable Execution

```
// data race
int a[] = {0,1};
std::vector<int> v;
for_each(par, a, a + 2, [&](int& i)
{
    v.push_back(i);
});
```

# Parallelizable Execution

```
// OK (don't do this):  
int a[] = {0,1};  
std::vector<int> v;  
  
std::mutex mut;  
for_each(par, a, a + 2, [&](int& i)  
{  
    mut.lock();  
    v.push_back(i);  
    mut.unlock();  
}));
```

# Parallelizable Execution

```
// OK (do this):  
int a[] = {0,1};  
std::vector<int> v(2);  
  
for_each(par, a, a + 2, [&](int& i)  
{  
    v[i] = i;  
}));
```

# Parallelizable Execution

```
// may deadlock (don't do this):
std::atomic<int> counter = 0;
int a[] = {0,1};

for_each(par, a, a + 2, [&](int& i)
{
    counter++;

    // spin wait for both lambdas to arrive
    while(counter.load() != 2)
    {
        ;
    }

    // try to do something crazy
    ...
});
```




# Parallelizable + Vectorizable Execution

```
algo(par_vec, begin, end, func);
```

`algo` is permitted to invoke function  
unsequenced if invoked in different threads,  
and unsequenced if invoked in the same thread

# Parallelizable + Vectorizable Execution

```
algo(par_vec, begin, end, func);
```



# Difference between `par` & `par_vec`

Function invocation is unsequenced when in different threads

When executed in the same thread

- `par`: unspecified, sequenced invocations
- `par_vec`: no sequence exists at all

# Parallelizable + Vectorizable Execution

It is the caller's responsibility to ensure correctness, for example that the invocation of functions do not attempt to synchronize with each other.

# Parallelizable + Vectorizable Execution

```
// may deadlock (don't do this):  
int counter = 0;  
int a[] = {0,1};  
  
std::mutex m;  
for_each(par_vec, a, a + 2, [&](int)  
{  
    mut.lock();  
    ++counter;  
    mut.unlock();  
}));
```

# Parallelizable + Vectorizable Execution

```
// OK:  
std::atomic<int> counter = 0;  
int a[] = {0,1};  
for_each(par_vec, a, a + 2, [&](int)  
{  
    ++counter;  
}));
```

# Parallelizable + Vectorizable Execution

```
// Best:  
int count = count_if(par_vec, ...);
```

Use the highest-level algorithm that does the job!

# Parallel Algorithms

Provide overloads of standard algorithm which receive execution policy as parameter

To the degree possible, parallelize everything



# New Algorithms

**reduce** : reduction over a collection

$$\text{result} = \text{init} + a[0] + a[1] + \dots + a[N-1]$$

**exclusive\_scan** : exclusive prefix sum

$$\text{result}[i] = \text{init} + a[0] + a[1] + \dots + a[i-1]$$

**inclusive\_scan** : inclusive prefix sum

$$\text{result}[i] = \text{init} + a[0] + a[1] + \dots + a[i]$$

# No Parallelism

Binary search algorithms

Heap algorithms

Permutation algorithms

Shuffling algorithms

Sequential numeric algorithms

# Parallel Exceptions

Parallel algorithms throw one of two:

- If no temporary storage is available,  
`bad_alloc`
- `exception_list` of exceptions thrown by  
user-provided code

# Exceptions Example

```
struct superstition_error { const char* what() { return "eek"; } };

std::vector<int> data = ...

try
{
    for_each(data.begin(), data.end(), [](auto x)
    {
        if(x == 13)
        {
            throw superstition_error();
        }
    });
}
catch(superstition_error& error)
{
    std::cerr << error.what() << std::endl;
}
```

# Exceptions Example

```
struct superstition_error { const char* what() { return "eek"; } };

std::vector<int> data = ...

using namespace std::experimental::parallelism;

try
{
    for_each(par, data.begin(), data.end(), [](auto x)
    {
        if(x == 13)
        {
            throw superstition_error();
        }
    });
}
catch(exception_list& error)
{
    std::cerr << "Encountered " << errors.size() << " unlucky numbers" << std::endl;

    reduce(par, errors.begin(), errors.end(), my_handler());
}
```

# Exceptions Example

```
struct superstition_error { const char* what() { return "eek"; } };

std::vector<int> data = ...

using namespace std::experimental::parallelism;

try
{
    for_each(seq, data.begin(), data.end(), [](auto x)
    {
        if(x == 13)
        {
            throw superstition_error();
        }
    });
}
catch(exception_list& error)
{
    std::cerr << "Encountered " << errors.size() << " unlucky numbers" << std::endl;

    reduce(seq, errors.begin(), errors.end(), my_handler());
}
```

# Example: Estimating $\pi$


# Example: Estimating Pi

Area of circle =  $\pi * \text{radius}^2$

Throw darts at unit square

$\pi \sim 4 * \text{green} / \text{total}$

Algorithm: Count the number of darts which fall within quarter circle





# Example: Estimating Pi

```
// generate a random point and test whether it
// lies in the quarter circle
bool test_quarter_circle(int seed)
{
    // seed an RNG
    std::default_random_engine rng(my_hash(seed));

    // generate numbers uniformly
    // in the unit interval
    std::uniform_real_distribution<float> u01(0,1);

    // generate a point within the unit square
    float x = u01(rng);
    float y = u01(rng);

    // measure distance from the origin
    float dist2 = std::sqrt(x*x + y*y);

    return dist2 <= 1.0f;
}
```

```
// burtleburtle.net/bob/hash/integer.html
int my_hash(int a)
{
    a = (a+0x7ed55d16) + (a<<12);
    a = (a^0xc761c23c) ^ (a>>19);
    a = (a+0x165667b1) + (a<<5);
    a = (a+0xd3a2646c) ^ (a<<9);
    a = (a+0xfd7046c5) + (a<<3);
    a = (a^0xb55a4f09) ^ (a>>16);
    return a;
}
```

# Example: Estimating Pi

```
// throw 300M darts
auto n = 300 << 20;

// create the integers [0..n)
auto iter = boost::make_counting_iterator(0);

// count the number of points in the quarter circle
using namespace std::experimental::parallel;

auto num_within_quarter_circle =
    count_if(par, iter, iter + n, test_quarter_circle);

double pi_estimate = (4.0 * num_within_quarter_circle) / n;
```

# Performance Portability

```
single CPU thread  
Intel i7 860
```

```
time ./pi_seq  
pi is approximately  
3.14
```

```
real 0m5.097s  
user 0m5.098s  
sys 0m0.004s
```

1x

```
8 CPU threads  
OpenMP, Intel i7 860
```

```
time ./pi_par  
pi is approximately  
3.14
```

```
real 0m0.971s  
user 0m7.565s  
sys 0m0.015s
```

5.25x

```
many GPU threads  
CUDA, NVIDIA Tesla K20
```

```
time ./pi_gpu  
pi is approximately  
3.14
```

```
real 0m0.260s  
user 0m0.047s  
sys 0m0.188s
```

19.6x

# Future Work

# Scheduling?

```
algo(exec, begin, end, func);
```

`algo` needs to compose with scheduling decisions in the surrounding application

`exec` specifies how `algo` is allowed to execute

- specifies **what** work an implementation is allowed to create
- does not specify **where** the work should be executed

Placement is **orthogonal**

# Scheduling?

```
algo(exec(sched), begin, end,  
func);
```

We anticipate extending our execution policies to accept scheduling requirements as parameters

`sched` specifies **where** the work should be executed

# Scheduling?

**sched** could be:

- hard requirement to execute vectorizable work in the current thread
- number of threads to use
- a thread pool to use
- an executor to use
- which GPU(s) to use

# Implementations in progress

[github.com/n3554](https://github.com/n3554)

- based on Thrust

[parallelstl.codeplex.com](https://parallelstl.codeplex.com)

- based on PPL?

[github.com/t-lutz/ParallelSTL](https://github.com/t-lutz/ParallelSTL)

- based on `std::thread`



# Summary

High-level algorithms make parallelism easy

- Portable & Composable
- Concurrency is invisible

Standardization

- On track for C++17
- Experimental Tech Spec in the meantime
- [github.com/cplusplus/parallelism-ts](https://github.com/cplusplus/parallelism-ts)

# Questions?

[jhoberock@nvidia.com](mailto:jhoberock@nvidia.com)  
[github.com/jaredhoberock](https://github.com/jaredhoberock)