

Rebuilding Boost Date Time for C++11/14

Jeff Garland
jeff@crystalclearsoftware.com

Alternate title:

C++11/14 Features for Building
Valuetype Classes - an exploration

What Will This Talk Achieve?

- Examination of C++11 & 14 features for building ValueType classes
 - Using boost date-time as an example
 - Differences from c++98
- General discussion of considerations for ValueType classes
- Preview of date-time v2

Background

- boost date-time (bdt)
 - v1 put into boost in 1.29 in 2002
 - used in many, many projects
- Good
 - Simple to use and fairly powerful interface
 - Fanatical error checking
 - I/O
- Bad
 - Fanatical error checking (it's somewhat slow)
 - Relatively hard to extend
 - Many code base hacks for pre-2000 compilers (eg: g++2.9.8, vc6)
- Ugly
 - I/O - facet based strftime based interfaces is relatively slow


```
//bdt v1
```

```
using namespace boost::gregorian;
```

```
date weekstart(2002, Feb, 1);    //obvious construction
```

```
//math
```

```
date weekend = weekstart + week(1);
```

```
date d2 = d1 + days(5);
```

```
//clock
```

```
date today = day_clock::local_day();
```

```
if (d2 >= today) {}    //all the usual comparison operators
```


date-time & chrono

- BDT v1 the basis is elements in <chrono>
- but...c++11 was NOT ready for a full date-time library
 - proposals were too late and immature (originally was targeting TR2)
 - date-time v1 is large enough that it's hard to standardize
 - committee was too busy
 - it was all we could do to get chrono...at the time
- Excuses :)

C++11/14 Features Covered

- Design considerations for the 'date' class (mostly)
- Specifically c++11/14 feature considerations (language and library)
 - final
 - noexcept (error handling)
 - move construction/assignment (R-values)
 - constexpr
 - member initialization
 - enum types
 - user defined literals
 - template aliases
 - std::to_string
 - delegating constructors
- Also, general class considerations
 - default construction
 - templated construction / conversion

Time to Start Over!

Reconsider everything...

What's easiest to use type in C++?

- int!
- Want date to be as easy as an int
- BDT v1 date is close...but not quite
 - as easy as int
 - or as fast
 - much closer to 'double'...which is a nightmare, really
- date can never be as good as int but we can do better than bdt v1

Properties of an Int

- Hard to break
 - Never throws exceptions
 - implicitly all functions are noexcept
 - used in virtually every program — mostly error free
 - fast at almost everything
 - now in C++11 good conversions from string
- What can go wrong with an int?
 - fail to initialize (easy fix there)
 - overflow on arithmetic
 - signed versus unsigned compare (warnings)

Design Considerations

- Fast (er)
- Simple (r) consistent interface
- I/O, I/O, I/O
- Extensible, extensible, extensible
- Play nice with <chrono>
- Play nice with BDT v1 — if possible...
- Something that can be standardized...

C++11/14 Features 1 by 1

final

- c++11 supports keyword 'final'
- In function context prevents derived classes from overloading
- In class context makes derivation an error

date as final class

```
class date final {.....};
```

```
class mydate : public boost::date_time2::date { }
```

```
> g++-4.9 -std=c++11 test.cpp
```

```
test.cpp:18:7: error: cannot derive from 'final' base  
'boost::date_time2::date' in derived type 'mydate'  
class mydate : public boost::date_time2::date
```


final - the final word

- Ultimately no classes in boost date_time2 are currently marked as final
- could prevent valid extension path — add stateless function / constructor
- no std library types use final


```
//bdt v1
```

```
using namespace boost::gregorian;
```

```
date weekstart(2002, Feb, 1);    //obvious construction
```

```
//math
```

```
date weekend = weekstart + week(1);
```

```
date d2 = d1 + days(5);
```

```
//clock
```

```
date today = day_clock::local_day();
```

```
if (d2 >= today) {}    //all the usual comparison operators
```


//input streaming

```
std::stringstream ss("2004-Jan-1");  
ss >> d3;
```

//date generator functions

```
date d4 = next_weekday(d3, Sunday);
```

//US labor day holiday is first Monday in Sept

```
nth_day_of_the_week_in_month labor_day(nth_dow::first, Monday, Sep);  
date d6 = labor_day.get_date(2004);
```


Do you see the problem?

- Too many ways to represent a date!
- And bdt v1 was inconsistent in its approach

It's all about construction

- default constructors
- move constructors (c++11)
- noexcept and errors (c++11)
- member initialization (c++11)
- templated constructor


```
date d3;  
d3 += days(2); //value?
```


It's all about construction

default constructor

- Should there be a default constructor?
 - initial answer was no in v1...
 - `std::map` requires key to be default constructible
 - some code clearly harder to write...
- What should it be?
 - current date...slow, slow, slow — breaks performance assumption
 - epoch - reasonable...answer for v2
 - 'not a date time'...answer for v1

you're so special - a diversion

- v1 had special values for date - nadt, neg_infinity, pos_infinity
- gone in v2
- advantages
 - handy for the domain on occasion
- disadvantages
 - requires addition of ~5 methods to date interface
 - date's aren't 'always valid' - can't reason about functions
 - checking for these - aka special logic
 - i/o is harder...
 - makes 'int' into 'double'

nadt == optional

- should be build optional into a low level value type?
- advantages
 - don't need the wrapper
- disadvantages
 - burdens all applications with optional behavior
 - even if not needed
 - most don't...
 - complexity...again


```
date d3;  
d3 += days(2); //v2 == epoch() + 2 days
```


It's all about construction

move construction?

- Should there be a move constructor?
 - value seems limited, however...
 - best practice to include anyway
 - can make it explicit in c++11

```
/// Trivial move constructor  
date(date&&) noexcept = default;
```

```
/// Trivial copy constructor  
date(const date&) noexcept = default;
```

```
/// Trivial assignment  
date& operator=(const date&) noexcept = default;
```


It's all about construction

Error Handling & NoExcept

- Problem: Sometimes you know your data is good...sometimes you don't
- Solution: checking versus non-checking constructor
- Non checking variant is no-except for max speed — use with trusted source
- Checking variant insures correctness

noexcept

```
date(const year_month_day& ymd) noexcept;  
date(const year_month_day& ymd, checking check);
```

- Performance boost for noexcept?
 - Unable to discern any at this point...
 - Lack of time & cleverness likely...

It's all about construction

member initialization

- Allows class/struct data members to be explicitly initialized


```
//user code...  
iso_week_number wn;
```

```
//library code...  
struct iso_week_number {  
    /** Construct to invalid state */  
    iso_week_number() noexcept = default;  
  
    ...  
  
    uint16_t year = 10001;  
    uint8_t week_no=54; ///use iso week numbering  
    uint8_t day_in_week=8; ///1==monday...7==sunday
```


Date - How do I represent thee?

- Strings — a gazillion variations
 - localization anyone?
 - iso
- month, day, year
- iso week number and day in week
- Calculated
 - third sunday in feb of 2014
 - last sunday in mar of 2014
 - sunday of week 5 in 2014
- modified julian day & julian day
- time_t and tm
- mayan calendar?
- klingon calendar?

It's all about construction

templated construction

- Really, I don't know how you're going to represent a date...seriously, I don't...

std::tm

```
//v1
tm d_tm;
d_tm.tm_year = 105;
d_tm.tm_mon = 0;
d_tm.tm_mday = 1;
date d = date_from_tm(d_tm);
```

```
//v2
tm t;
t.tm_year = 113;//2013
t.tm_mon = 11; //Dec
t.tm_mday = 30;
date d(t);
```


std::tuple & chrono

//v2 only — can't write in v1

```
std::tuple<int, int, int> t_ymd(1900, 1, 1);  
date d(t_ymd);
```

```
using namespace std::chrono;  
system_clock::time_point tp = system_clock::now();  
date d(tp); //v2 only
```



```
//v2 api
```

```
using namespace boost::date_time2;
```

```
date d(2012, 1, 1);
```

```
year_month_day ymd(2012, 1, 1);
```

```
date d2(ymd);
```

```
date d1(year_month_day("20140401T000000"), ISO);
```

```
date d2(year_month_day("2014-04-01"));
```

```
date d1(year_month_day("2004-10-01"));
```

```
date d2(year_month_day("2004/10/01"));
```

```
//calculated dates
```

```
day_of_week dow(First, Wed, Jan, 2013 );
```

```
date d(dow);
```

```
date d(day_of_year(2014, 1)); //jan 1, 2014
```

```
closest_day_of_week pdw(Sun, Before, 2013, May, 17);
```

```
date d(pdw);
```



```
iso_week_number wn1("2014-W01-2");  
date d1(wn1);
```

```
iso_week_number wn2("2014", "W1", "2");  
date d2(wn2);
```

```
iso_week_number wn3(2014, 1, 5);  
date d3(wn3); //2014-Jan-3
```

```
//bridge from bdt v1  
boost::gregorian::date bd(1900, 1, 1);  
date d(bd);
```

```
boost::posix_time::ptime pt(bd);  
date d2(bd);
```


How is it done?

- Templated constructor
- Specializations for different types
- Allows users to add new representation — all construction is the same
- Alternative
 - Construct everything using `make_date` function
 - Users can provide their own
 - Feels odd for `date`...

Under the Hood

```
//declaration...
template<typename T>
explicit date(const T& t) noexcept;

/// Template specialization to construct a date from a chrono::system_clock::time_point
template<>
date::date<std::chrono::system_clock::time_point>(const std::chrono::system_clock::time_point& tp) noexcept
{
    using namespace std::chrono;
    std::time_t tt = system_clock::to_time_t(tp);
    from_time_t(tt);
}
```


User Defined Literals

- C+11 allows creation of user defined literal
- Convert a literal in code to a type
- C++14 has pre-defined values for chrono

user defined literal example

```
#include <chrono>
#include <iostream>
using namespace std::chrono;
using namespace std::literals::chrono_literals;

hours h(1); //traditional
auto ns = 1h + 20us; //ns type == chrono::nanoseconds
std::cout << ns.count() << std::endl;
```


Nice, but...

- Standard limits user defined literals
 - must include underscore
 - `_w` for week?
- Types are limited...
- Easy to implement?

snippet of chrono 'h' operator

```
constexpr chrono::duration<long double, ratio<3600,1>>  
operator""h(long double __hours)  
{ return chrono::duration<long double, ratio<3600,1>>{__hours}; }
```

```
template <char... _Digits>  
constexpr typename  
__select_type::_Select_type<__select_int::_Select_int<_Digits...>::value,  
                           chrono::hours>::type  
operator""h()  
{  
    return __select_type::_Select_type<  
        __select_int::_Select_int<_Digits...>::value,  
        chrono::hours>::value;  
}
```


what's wrong with this?

```
constexpr boost::date_time2::days operator""_d(short d)
{
    return boost::date_time2::days(d);
}
```

```
g++-4.9 -I ../bdt2 -std=c++14 test.cpp
```

```
test.cpp:22:45: error: 'boost::gregorian::days boost::date_time2::literals::operator""_d(short int)' has invalid
argument list
```

```
    boost::date_time2::days operator""_d(short d)
                                   ^
```


still no joy...

```
constexpr days operator""_d(unsigned long long d)
{
    return days(d);
}
```

error...

boost::date_time2::days' is not an aggregate, does not have a trivial default constructor, and has no constexpr constructor that is not a copy or move constructor

Remove the constexpr for now

```
constexpr days operator""_d(unsigned long long d)
{
    return days(d);
}
```

- question: is it worth it to be able to write?

```
auto d = 1_d + 2_w; //15 days
```

```
date d {2014, 1, 1};
```

```
d += 1_w;
```

- ???

constexpr

- Generalized constant expression
- Function evaluated at compile time
- Some C++11 limits
 - typically needs a single return value
 - can only call other constexpr functions
- C++14 generalizes constexpr
 - Allows control structures (if/switch)
 - Local variables

obvious constexpr

```
static constexpr uint8_t day_of_month_min() { return 1; }  
static constexpr uint8_t day_of_month_max() { return 31; }  
  
static void validate_ymd(uint16_t year, uint8_t month, uint8_t day)  
{  
    range_check("year", year_min(), year_max(), year);  
    range_check("month", month_min(), month_max(), month);  
    range_check("day", day_of_month_min(), day_of_month_max(), day);  
}
```


constexpr limits

- consider 'max_date' function
- returns a date that represents max representable date

date::max_date

```
static constexpr date max_date()
{
    return date(year_max()-1,
                month_max(),
                day_of_month_max());
}
```


well, maybe not...

```
g++-4.8 -I ~/devtools/boost_1_55_0 -std=c++11 test.cpp
```

```
In file included from test.cpp:6:0:
```

```
date.hpp: In static member function 'static constexpr boost::date_time2::date
```

```
boost::date_time2::date::max_date()':
```

```
date.hpp:85:29: error: invalid return type 'boost::date_time2::date' of constexpr function 'static constexpr
```

```
boost::date_time2::date boost::date_time2::date::max_date()'
```

```
    static constexpr date max_date()
```

```
        ^
```

```
date.hpp:78:11: note: 'boost::date_time2::date' is not literal because:
```

```
    class date : public date_base<gregorian_calendar>
```

```
        ^
```

```
date.hpp:78:11: note: 'boost::date_time2::date' is not an aggregate, does not have a trivial default constructor, and has no constexpr constructor that is not a copy or move constructor
```


std::to_string

- new standard library functions to convert many integral types to std::string
- types covered include integer and floating point of various flavors
- corresponding string to type (eg: stoi) functions also there

using to_string - generic range check

```
template <typename T>
void range_check(std::string unit, T min, T max, T value)
{
    if (value > max || value < min) {
        throw std::out_of_range(unit + " is out of range "
            + std::to_string(min) + "..." + std::to_string(max)
            + ": " + std::to_string(value));
    }
}
```


delegating constructors

- Allows calling one constructor from another
- Avoids writing 'init' type functions

delegating constructor

```
class year_month_day {  
public:  
    year_month_day(const char* const ymd_string);  
  
    template<typename T>  
    explicit year_month_day(const T& ymd);  
};
```

```
template<>  
year_month_day::year_month_day(const std::string& ymd_string) :  
    year_month_day(ymd_string.c_str())  
{  
}
```


Building Valuetypes in C++ +11/14

- It's a whole new world...
- Top features — Jeff's view...
 - to_string
 - noexcept
 - explicit defaults for compiler generated constructors
 - member initialization
 - constexpr
 - delegating constructor
- questionable value
 - final
 - user defined literals

Conclusions: Valuetype Design Considerations

- How does your type integrate with others?
- Does it play well with things in standard library?
- Does it follow recognized / common patterns from standard?

Conclusion

- C++11 and 14 provide nice features for writing value types
- Implementation quality is much higher
- But don't forget about old features!
- Other resources
 - Sean Parent - C++Now 2012 - Value Semantics and Concepts-Based Polymorphism
 - Eric Niebler C++ Now 2014 - C++11 Library Design
 - Michael Caisse - The Canonical Class - Wed 9 am